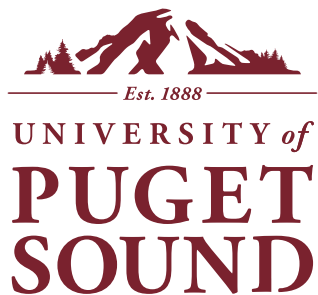


CS 455

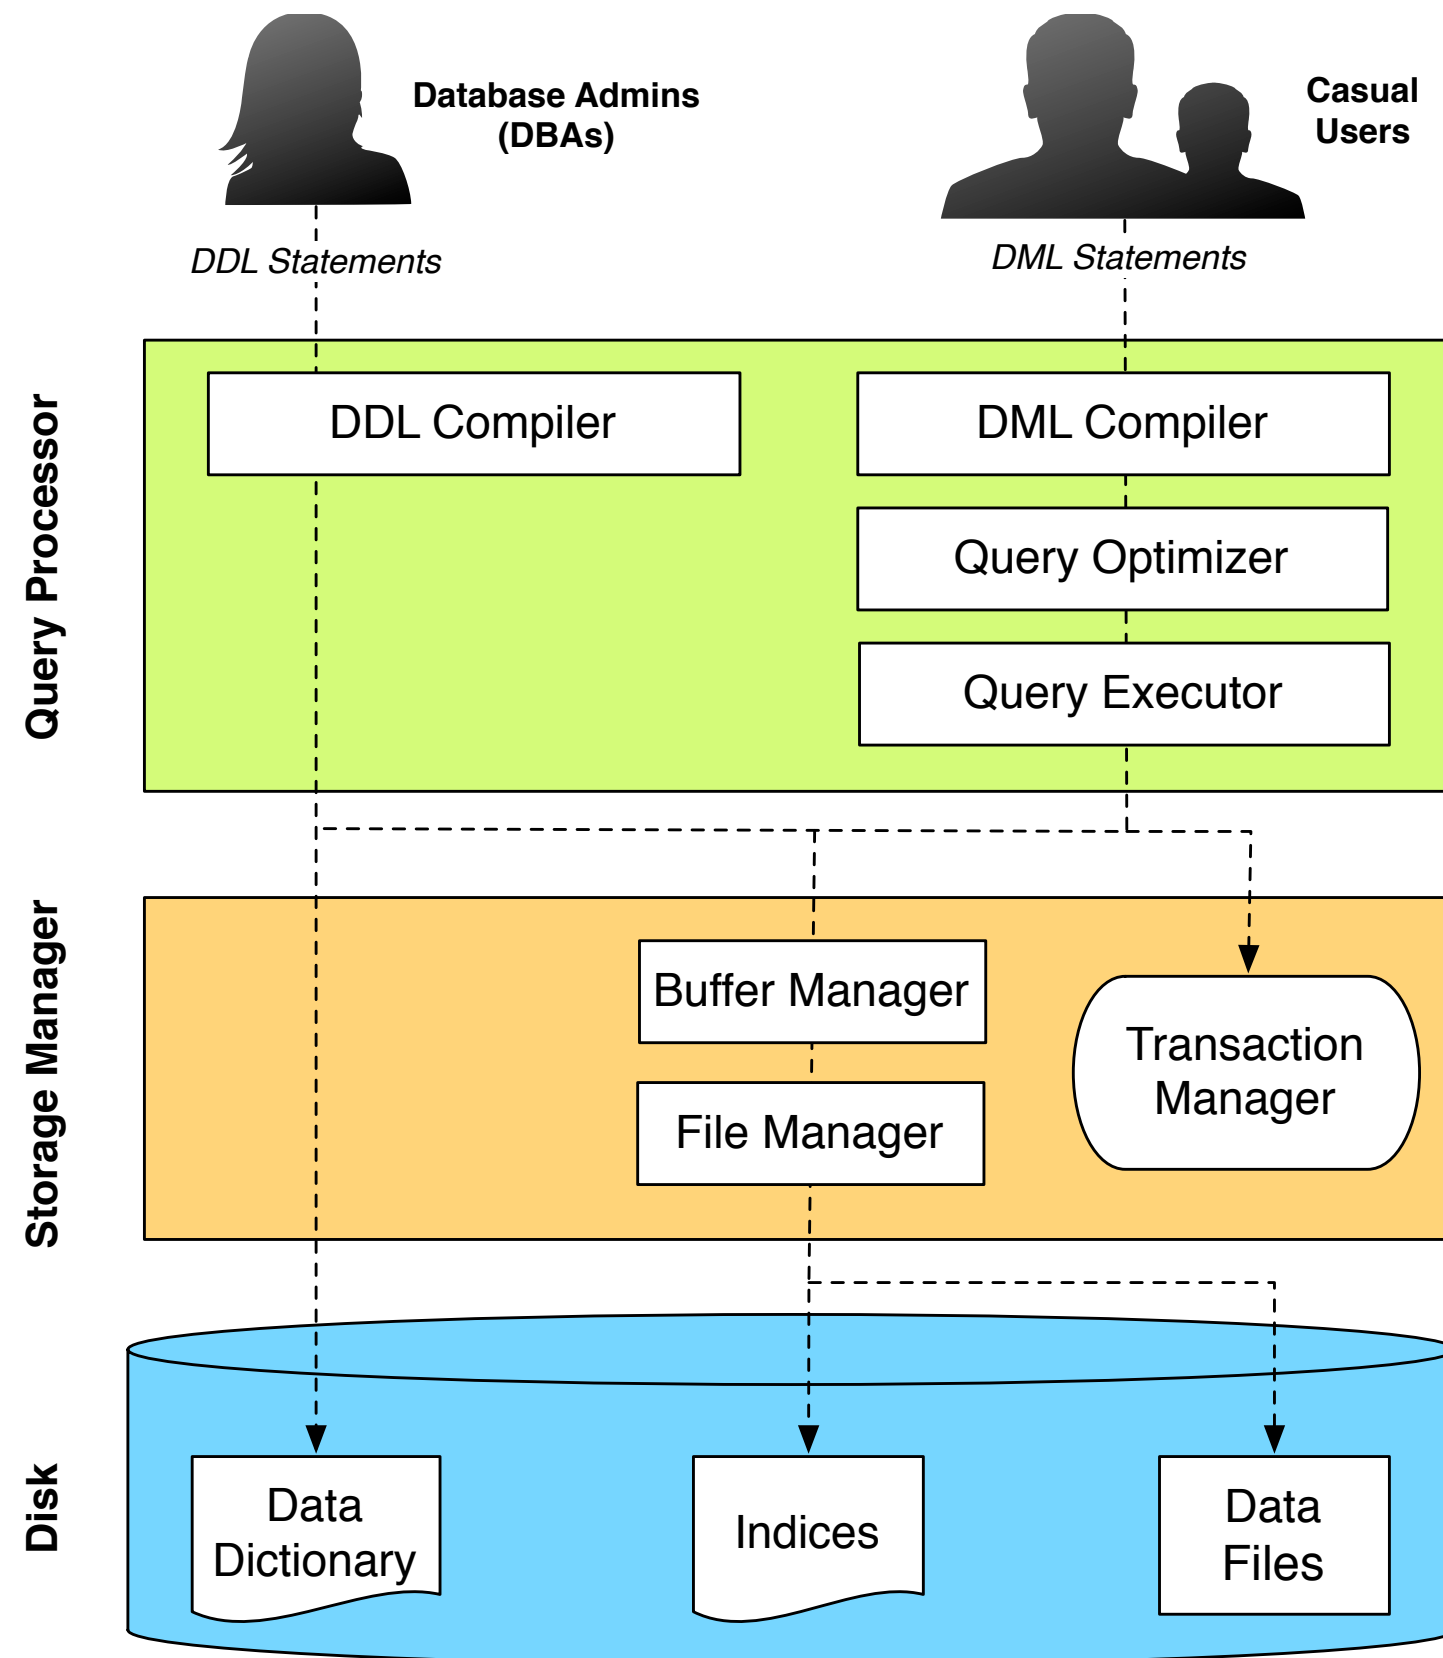
Principles of Database Systems



Department of Mathematics
and Computer Science

Lecture 5
Persistent Storage and Files

Relational Database System Architecture



Topics

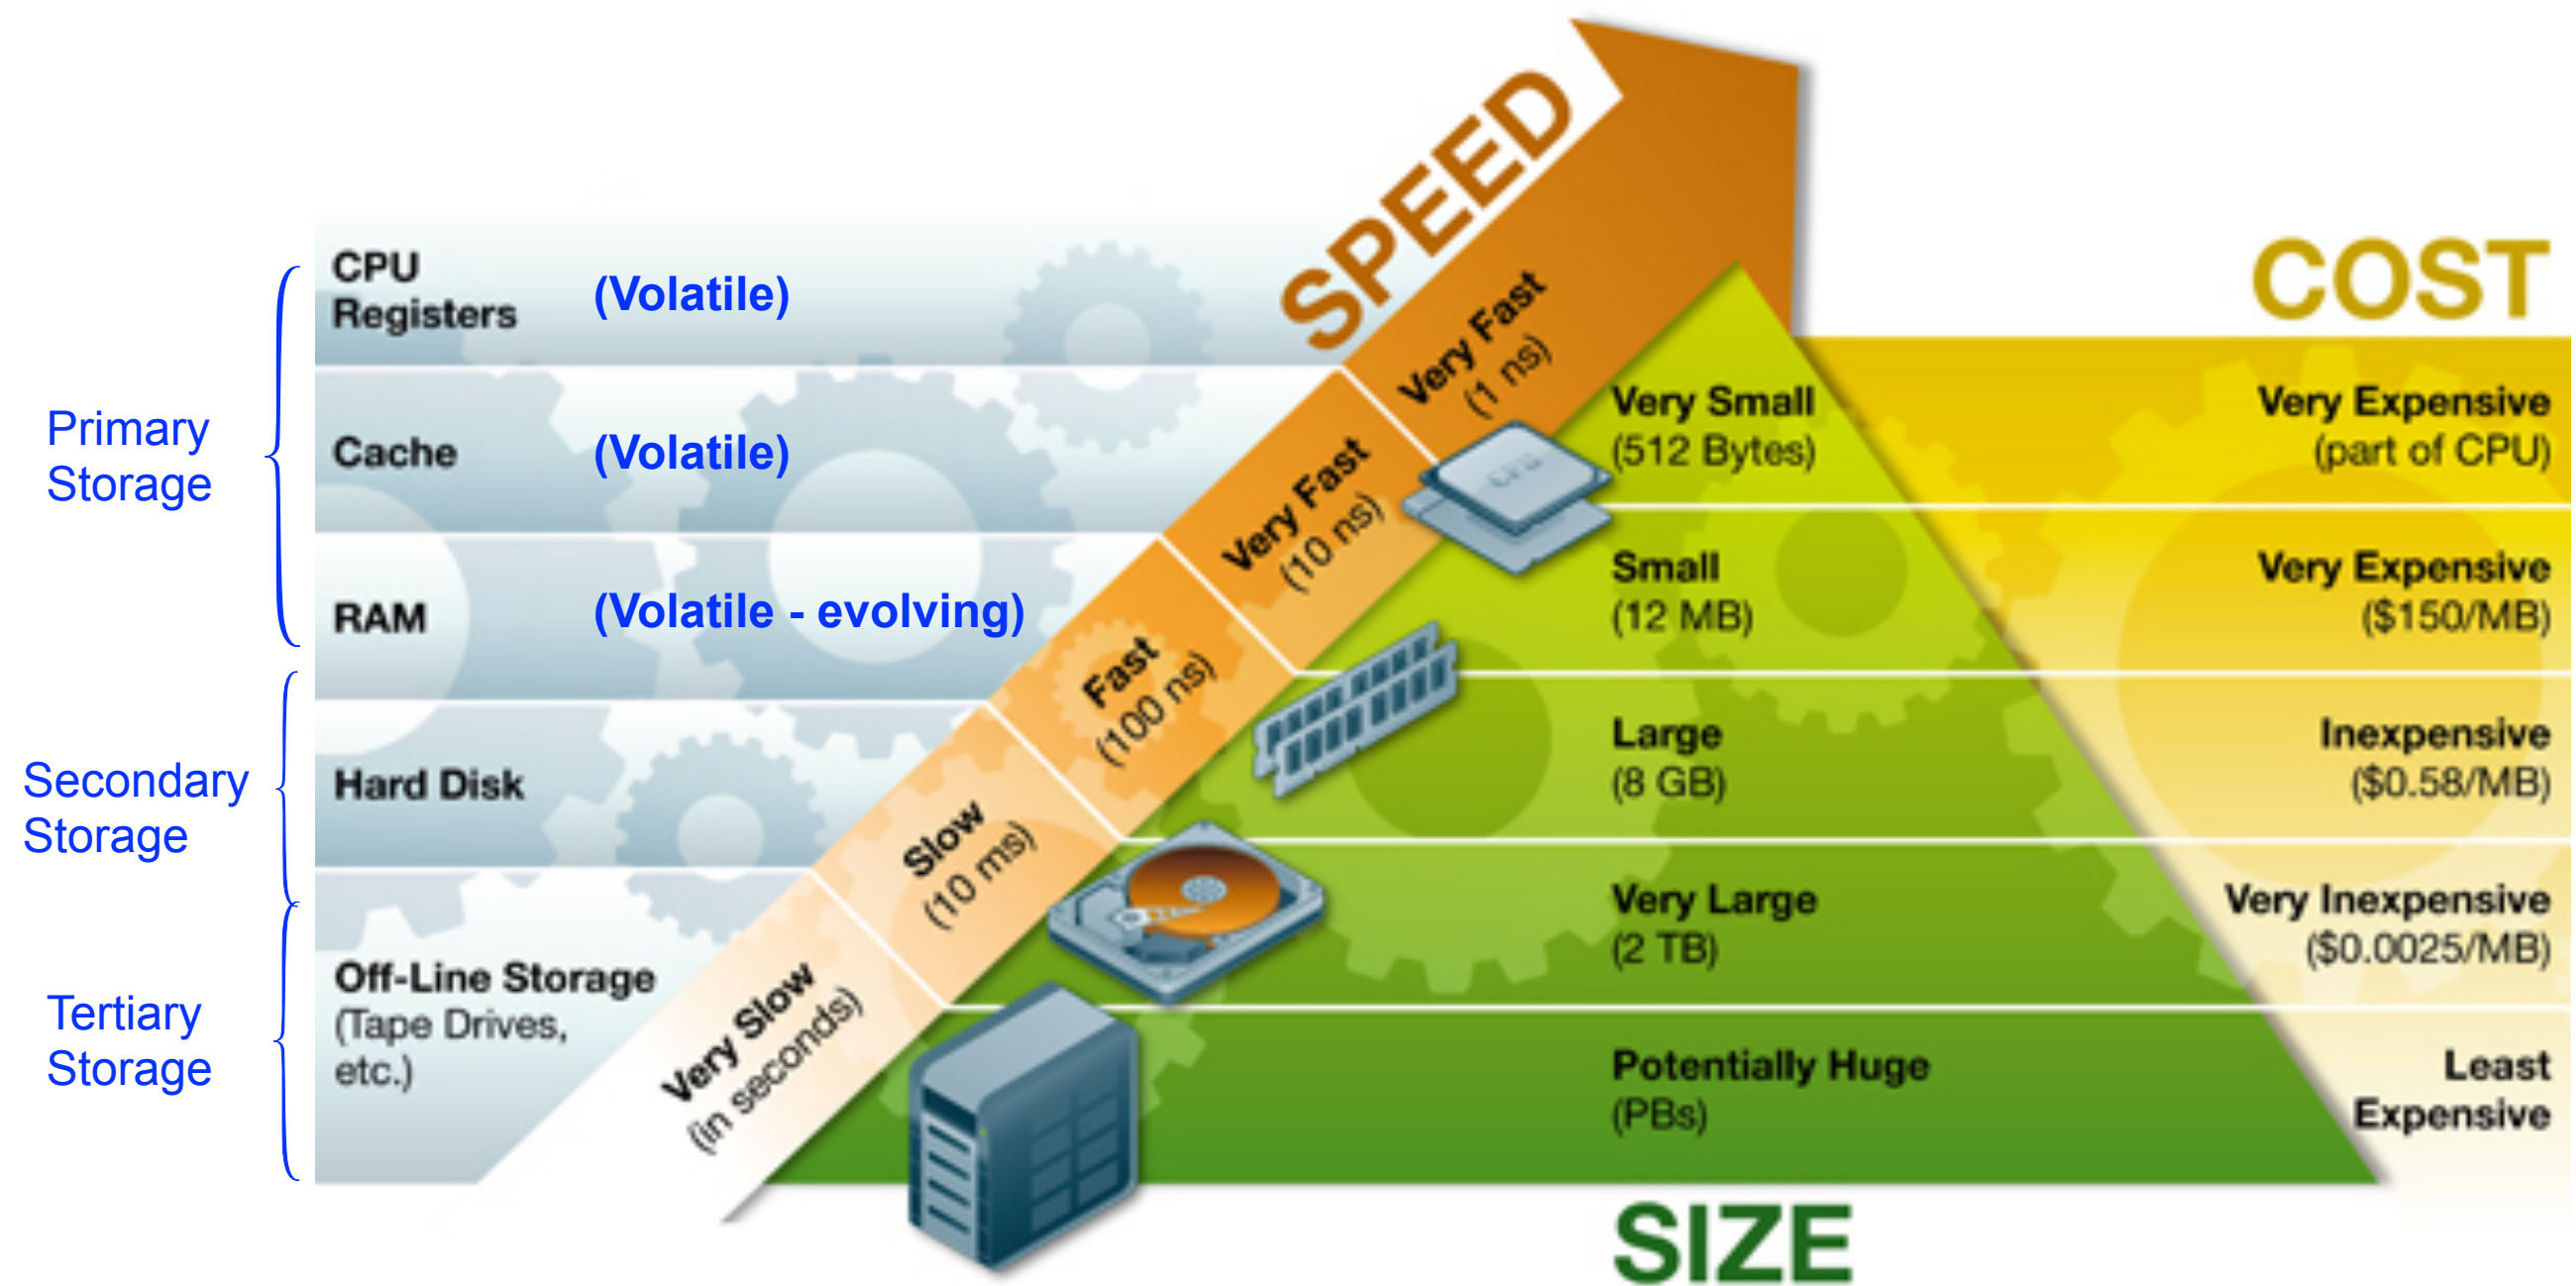
► Disks

- Hardware
- Understanding Files and Blocks
- Modeling Access Time
- Disk Scheduling

► Buffer Management



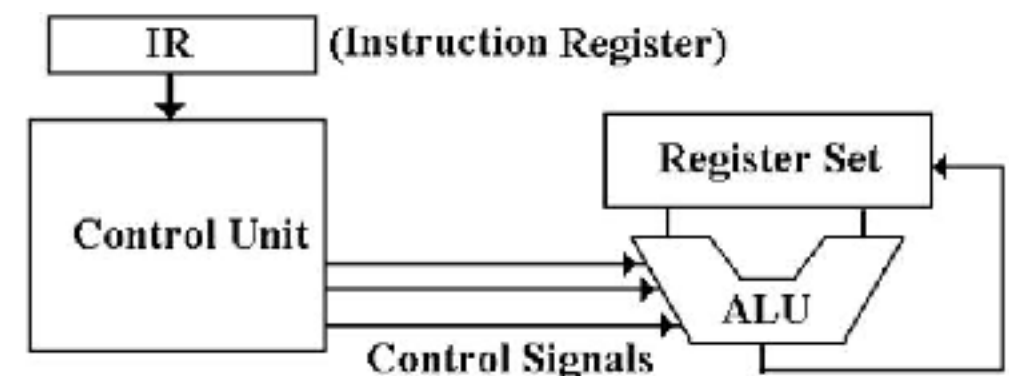
Storage (Memory) Hierarchy



Storage Hierarchy: CPU Registers

► CPU Registers (Volatile!)


- Fastest level of storage and the closest to computation
 - The CPU's arithmetic logic unit (ALU) performs calculations and comparisons directly off register data
- Each register stores 32 or 64 bits (of instruction or data)
 - Only a few registers, even on modern CPUs
 - ~32 integer registers, 32 floating point registers for MIPS CPUs
 - ~8 on x86 (Intel CPUs)
 - Why not pack more?




Storage Hierarchy: Cache

► CPU Cache (Volatile!)

- A few KBs to several MBs
- Resides on CPU
 - Modern CPUs have several “levels” of cache before exploring memory
 - Levels: L1, L2, and L3 are common
 - Sizes: $L1 < L2 < L3$
- Automatically managed by hardware
 - We can't program caches to do what we want... but we can be better programmers



Intel® Core™ i7-3960X Processor Die Detail



Storage Hierarchy: Main Memory

► Main Memory (Volatile!)

- Several GBs
 - New iMacs: 32 GB RAM
- Still fast(-ish): 10-100 nanoseconds (ns)
- Still quite costly
 - \$0.0054 per MB

<http://www.jcmit.com/memoryprice.htm>



- ## ► In general, still too small and expensive to hold DBs
- Also, no "data persistence" at this level

Storage Hierarchy: Hard Drives

► Hard Drives (Non-Volatile!)

- A few TBs
 - New 2013 iMacs: 1 TB
 - Our HPC cluster node: 4 TBs (Two 2 TB disks)
- Very cheap
 - \$0.0000467 per MB

<http://www.jcmit.com/diskprice.htm>



► A DB is stored here, but...

- Hard drives are slow: Few milliseconds (ms)
 - Why so slow?
 - We need to understand how it's built

Storage Hierarchy: Solid-State Disks

► Solid State Disks or SSDs (Nonvolatile!)

- Hundreds of GBs to a few TBs
- Faster than magnetic disks: 100 microseconds (us)
- An order of magnitude more expensive than magnetic drives
 - \$0.000518 per MB

<http://www.jcmit.com/flashprice.htm>

► *Is write endurance a problem?*

- A block of memory can endure 2000 to 3000 write operations



"This breed of non-volatile storage retains data by trapping electrons inside of nanoscale memory cells. A process called tunneling is used to move electrons in and out of the cells, but the back-and-forth traffic erodes the physical structure of the cell, leading to breaches that can render it useless."

<http://techreport.com/review/27909/the-ssd-endurance-experiment-theyre-all-dead>

Storage Hierarchy: Offline Storage

► Offline (Tertiary) Storage (NonVolatile)

- Optical drives, tapes, removable flash, cloud storage services (S3, DropBox)
- Enormous amounts of offline storage
- Extremely cheap and expendable



► Way too slow for production DB applications

- But are still relevant today
- Mostly for backup and portability



Today: Magnetic Disks Still Dominate

► First modern magnetic HDD with moving head

- Shipped with IBM 305 RAMAC (1956)
- Held 5 MB for a whopping \$10,000 per MB!
- 24 inches diameter, single-head



► Newest HDD today (2021)



- 10 TB, 7200RPM, 3.5 inches under \$300 total < \$0.0000003 per MB

► Full History


- <http://www.pcworld.com/article/127105/article.html>

Magnetic Disk Growth

- Disks over the years:
 - Exponentially dense
 - Exponentially cheaper per MB stored
 - But what about disk performance?



Magnetic Disk Growth (Cont.)




Credits: Jim Gray. "Rules of Thumb in Data Engineering."

- ▶ Disk capacity grows 60% per year
- ▶ Disk performance grows 10 times slower!
 - 7% per year
- ▶ Legend:
 - tpi = tracks per sq-inch
 - kbpi = kilobits per sq-inch
 - MBps = MB/sec

Hard Disk Drives (HDD)

- ▶ A disk has of one or more *platters*
 - Each *platter* has two *surfaces*
 - Each *surface* has many *tracks*
 - 2 to 300,000 tracks per surface
 - Each *track* has many *sectors*
 - Each *sectors* stores **512 Bytes**
- ▶ Each *platter* has a *head* that performs read/write operations
- ▶ *Cylinders* are the set of the same tracks on all surfaces



Communication Media

► Disks are connected to the host machine via a communication medium

► Cables are rated by *transfer rates (or bandwidth)*:

- IDE (ATA): ~2.8 Gbps
- SATA 3 (Serial ATA): 6.0 Gbps
- SCSI (pronounced “scuzzy”): 5.0 Gbps
- USB 3.0: Originally 5.0 Gbps
 - (Now 10 Gbps to compete with Apple Thunderbolt)



Topics

► Disks


- Hardware
- Understanding Files and Blocks
- Modeling Access Time
- Disk Scheduling

► Buffer Management




Disk Blocks

- ▶ **Block:** Contiguous sequence of sectors from a single track
 - Data is read/written a block at a time
 - Standard block size is usually 4 KB
 - (8 * 512 Byte sectors)
- ▶ **Important!** Assume **blocks** are the basic unit of data transfer
 - Not sector
 - Even though "sector size" often synonymous with "block size"




Translating User to System View

- ▶ What happens if user says, "**read** bytes 2 - 14 from a file?"
 - OS will fetch disk block(s) corresponding to those bytes
 - Return just the correct portion of the block(s) to CPU
 - Assume bytes 2-14 in the file is in **Block 4** on disk



Translating User to System View

- ▶ What happens if user says, "**write** bytes 2 - 14 into a file?"
 - OS will fetch disk block(s) corresponding to those bytes
 - Modify portion in memory
 - Write out block to disk:
 - Accesses increases by factor of 2! Writes are expensive!




Then What's a File?

- ▶ A *File* is just an operating system abstraction for a sequential stream of bytes.





Then What's a File?

- ▶ A *File* is just an operating system abstraction for a sequential *stream* of bytes.
- ▶ To us: We *think* that a file is just one contiguous unit kept together on disk




Then What's a File?

- ▶ A *File* is just an operating system abstraction for a sequential *stream* of bytes.
- ▶ To us: We *think* that a file is just one contiguous unit kept together on disk
- ▶ In reality: Every file is split into blocks
 - Blocks themselves may or may **NOT** be arranged contiguously on disk
 - If we're **unlucky**, file blocks are scattered across disk (file is "**fragmented**")




Then What's a File?

- ▶ A *File* is just an operating system abstraction for a sequential *stream* of bytes.
- ▶ To us: We *think* that a file is just one contiguous unit kept together on disk
- ▶ In reality: Every file is split into blocks
 - Blocks themselves may or may **NOT** be arranged contiguously on disk
 - If we're *lucky*, the file blocks are also contiguous on disk!



The File Illusion


- ▶ Example: block size = 4 KB, and some file A.txt is 11 KB
 - A.txt requires $\text{ceiling}(11 \text{ KB} / 4 \text{ KB}) = 3$ blocks
 - OS allocates these disk blocks: 4, 0, 5



The File Illusion

- ▶ Example: block size = 4 KB, and some file A.txt is 11 KB
 - A.txt requires $\text{ceiling}(11 \text{ KB} / 4 \text{ KB}) = 3$ blocks
 - OS allocates these disk blocks: 4, 0, 5
 - User thinks the 3 blocks are contiguous
 - These are called logical blocks: 0,1,2


```
Scanner file = new Scanner(new File("A.txt"));
while (file.hasNext()) {
    // we think the file data is contiguous!
    // give me the next line/row of the file!
    String s = file.nextLine();
}
```



Files are *usually* read sequentially (linear fashion)


Data Access Patterns Matter

- ▶ Assume a file's blocks can be arranged contiguously on disk
 - The way we access files (known as *access patterns*) still matter
- ▶ Access Pattern:
 - How is a file read from (or written to) disk?
 - Sequential access pattern: [1,2,3,4,5,6,7,8]




Data Access Patterns Matter

- ▶ Assume a file's blocks can be arranged contiguously on disk
 - The way we access files (known as *access patterns*) still matter
- ▶ Access Pattern:
 - How is a file read from (or written to) disk?
 - Sequential access pattern: [1,2,3,4,5,6,7,8]




Data Access Patterns Matter

- ▶ Assume a file's blocks can be arranged contiguously on disk
 - The way we access files (known as *access patterns*) still matter
- ▶ Access Pattern:
 - How is a file read from (or written to) disk?
 - *Sequential access pattern: [1,2,3,4,5,6,7,8]*
 - e.g., scanning through the file line-by-line
 - e.g., linear search over a huge array stored in file
 - Ideal situation for performance!




Data Access Patterns Matter

- ▶ Assume a file's blocks can be arranged contiguously on disk
 - The way we access files (known as *access patterns*) still matter
- ▶ Access Pattern:
 - How is a file read from (or written to) disk?
 - Random (irregular) access pattern [6,1,8,3,2,4,7,5]




Data Access Patterns Matter

- ▶ Assume a file's blocks can be arranged contiguously on disk
 - The way we access files (known as *access patterns*) still matter
- ▶ Access Pattern:
 - How is a file read from (or written to) disk?
 - Random (irregular) access pattern [6,1,8,3,2,4,7,5]
 - Approximates file fragmentation again!
 - Random access is bad for performance!



Data Access Patterns Matter

- ▶ Assume a file's blocks can be arranged contiguously on disk
 - The way we access files (known as *access patterns*) still matter
- ▶ Access Pattern:
 - How is a file read from (or written to) disk?
 - Random (irregular) access pattern [6,1,8,3,2,4,7,5]
 - Random access is bad for performance!
 - Approximates file fragmentation again!



Main Takeaways

- ▶ Files appear like a sequential entity to humans, but that's an *illusion*
 - Your OS splits the file into blocks
 - Finds a free block on disk to store each chunk
 - Sequential ordering of blocks for a file is **not** guaranteed
 - File may be "fragmented" across disk
 - What does fragmentation do to performance of reading the file?

Topics

► Disks

- Hardware
- Understanding Files and Blocks
- **Modeling Access Time**
- Disk Scheduling

► Buffer Management



Modeling Average Data Access Time (DAT)

$$DAT = T_{decode} + T_{seek} + T_{lat} + T_{transfer}$$

- Decoding Overhead (T_{decode})
 - Disk controller decodes the R/W command, initializes physical movements
- Seek Time (T_{seek})
 - Time taken to move arm to the desired track.
 - **Moving to adjacent tracks = 0 cost.**
- Rotational Latency (T_{lat})
 - Spin the desired data under the arm. Reciprocal of RPM (rotations per minute)
 - **On average**, you need **(full cycle/2)** rotations per block if block are not contiguous
- Transfer Time ($T_{transfer}$)
 - Time taken to send data payload to/from host machine

Example

$$DAT = T_{decode} + T_{seek} + T_{lat} + T_{transfer}$$

- ▶ Estimate the amount of time it would take to read a 1 GB file off a disk with the following characteristics
 - Sector (and block) size = 4 KB; Track size = 2 GB
 - Decode overhead = 0.1 ms per block fetched
 - 7200 rotations/minute (RPM)
 - 4ms seek time per sector (if next block is on a different track)
 - Disk is connected to host via USB 3.0 (= 5/8 GB/s)
 - Assume sequential access pattern
- ▶ File stored contiguously vs File stored randomly on disk (fragmented)

Seek Time + Latency Dominate

$$DAT = T_{decode} + T_{seek} + T_{lat} + T_{transfer}$$

- ▶ If file is **contiguously** arranged on disk and access is **sequential**
 - $DAT = 26214.4\text{ms} + 4\text{ms} + 4.167\text{ms} + 1600\text{ms}$
 - $DAT = 27.82 \text{ sec}$
- ▶ If file is **fragmented** on disk and/or access is **irregular**
 - $DAT = 26214.4\text{ms} + 2^{20}\text{ms} + 2^{18} \cdot 4.167\text{ms} + 1600\text{ms}$
 - $DAT = 2168.7 \text{ sec} = \sim 36 \text{ minutes}$

Accelerating (Magnetic) Disk-Access Times

- ▶ Try to keep blocks pertaining to same file close together on disk
 - Contiguously if possible (to accelerate sequential file access)
 - *Whose job? Operating System (file system)*
- ▶ Access disk data in a sequential pattern
 - Random access patterns approximate block-level file fragmentation
 - e.g., `nextLine()` is good for performance
 - *Whose job? Programmer*

Accelerating (Magnetic) Disk-Access Times (2)

- ▶ Speed up disk rotations
 - Can we reduce latency?
 - *Whose job? Computer engineers*
 - *(But 7200 RPM is already near physical limitations)*

- ▶ Dispatch (schedule) disk operations in a smarter way
 - Can we reduce seek time?
 - *Whose job? DB and/or OS (Next)*

Topics

► Disks

- Hardware
- Understanding Files and Blocks
- Modeling Access Time
- **Disk Scheduling**

► Buffer Management



Copying a File



Disk Scheduling

► **Problem:** *Seek time is a dominant term in DAT.*

► *Disk Scheduling*

- Given multiple outstanding disk block requests, in what order should we dispatch them to the disk?

► **Goals:** We want a schedule that

- Minimizes total latency (delay) for *small* transfers
 - e.g., a selective database query that only returns a couple records from a table
- Maximizes throughput (bytes per unit time) for *large* transfers
 - e.g., returning all records in a large database table

FCFS Policy

First-Come-First-Served Policy

- Schedule each disk I/O operation in the order it arrives

► Pros

- It's fair! Requests honored in the order in which they arrive
- Easy to manage pending requests with a queue (fast!)

► Cons


- Disk-head locality is not taken into consideration
- Wide alternating arm swings are possible.

FCFS Example

- ▶ Disk head currently on track **100**
- ▶ Outstanding requests for I/O to blocks on tracks:
 - **190**, **55**, **190**, **74**, **40**, **190**
 - (Different colors = concurrent DB transactions)
- ▶ Let's define a simple metric to compare policies
 - *Average Seek Length per Block*
 - *(Tracks traversed / number of blocks fetched) in the schedule*
 - *Approximates seek time*

FCFS Example

- ▶ Disk head currently on track **100**
- ▶ Outstanding requests for I/O to blocks on tracks:
 - **190**, **55**, **190**, **74**, **40**, **190**
 - (Different colors = concurrent DB transactions)



SSTF Policy

- Problem with FCFS is that disk-head location is not considered

Shortest-Seek-Time-First (SSTF) Policy

1. Keep track of the current head position
2. Always issue the I/O operation that requests a track nearest to the head
3. If there's a tie, use FCFS or flip a coin

SSTF Example


- ▶ Disk head currently on track **100**
- ▶ Outstanding requests for I/O to blocks on tracks:
 - **190**, **55**, **190**, **74**, **40**, **190**
 - (Same example as before)
- ▶ Evaluation
 - Average seek length = ?
 - *Is SSTF **optimal** for minimizing average seek length?*
 - Potential problems?

SSTF Example

- ▶ Disk head currently on track **100**
- ▶ Outstanding requests for I/O to blocks on tracks:
 - **190**, **55**, **190**, **74**, **40**, **190**
 - (Same example as before)
- ▶ Evaluation
 - Average seek length = ?
 - Is SSTF *optimal* for minimizing average seek length?
 - Potential problems?
 - *Starvation* is possible for outer tracks as requests continually arrive? (Why?)


LOOK Policy (Also called Elevator)

- ▶ Starvation is a problem!




LOOK Policy

1. Disk-head has *direction* and *current_track*
2. Sort requests in order of track number according to *direction*
3. If no more I/Os to service in *direction*, reverse direction




LOOK Policy




LOOK Policy

1. Disk-head has *direction* and *current_track*
2. Sort I/O in order of track number according to *direction*
3. If no more I/Os to service in *direction*, reverse direction




LOOK Policy




LOOK Policy

1. Disk-head has *direction* and *current_track*
2. Sort I/O in order of track number according to *direction*
3. If no more I/Os to service in *direction*, reverse direction




LOOK Policy




LOOK Policy

1. Disk-head has *direction* and *current_track*
2. Sort I/O in order of track number according to *direction*
3. If no more I/Os to service in *direction*, reverse direction




LOOK Policy




LOOK Policy

1. Disk-head has *direction* and *current_track*
2. Sort I/O in order of track number according to *direction*
3. If no more I/Os to service in *direction*, reverse direction




LOOK Policy




LOOK Policy

1. Disk-head has *direction* and *current_track*
2. Sort I/O in order of track number according to *direction*
3. If no more I/Os to service in *direction*, reverse direction




LOOK Policy




LOOK Policy

1. Disk-head has *direction* and *current_track*
2. Sort I/O in order of track number according to *direction*
3. If no more I/Os to service in *direction*, reverse direction




LOOK Policy



LOOK Policy

1. Disk-head has *direction* and *current_track*
2. Sort I/O in order of track number according to *direction*
3. If no more I/Os to service in *direction*, reverse direction



LOOK Example

- ▶ Disk head currently on track **53**
- ▶ Outstanding requests for I/O to sectors on cylinders:
 - 98, 183, 37, 122, 14, 124
- ▶ Evaluation
 - Average seek length = ?
 - Starvation reduced, but still possible

Circular LOOK Policy

- ▶ How to guarantee all tracks are visited periodically?
 - Think typewriter
 - (Not implemented in practice -- too wasteful)



Circular LOOK Policy

1. Just like LOOK, but head moves in only one *direction*
2. If no more I/Os to service in *direction*, move head to track #0
(Ignoring any requests on the way back to track #0)

Topics

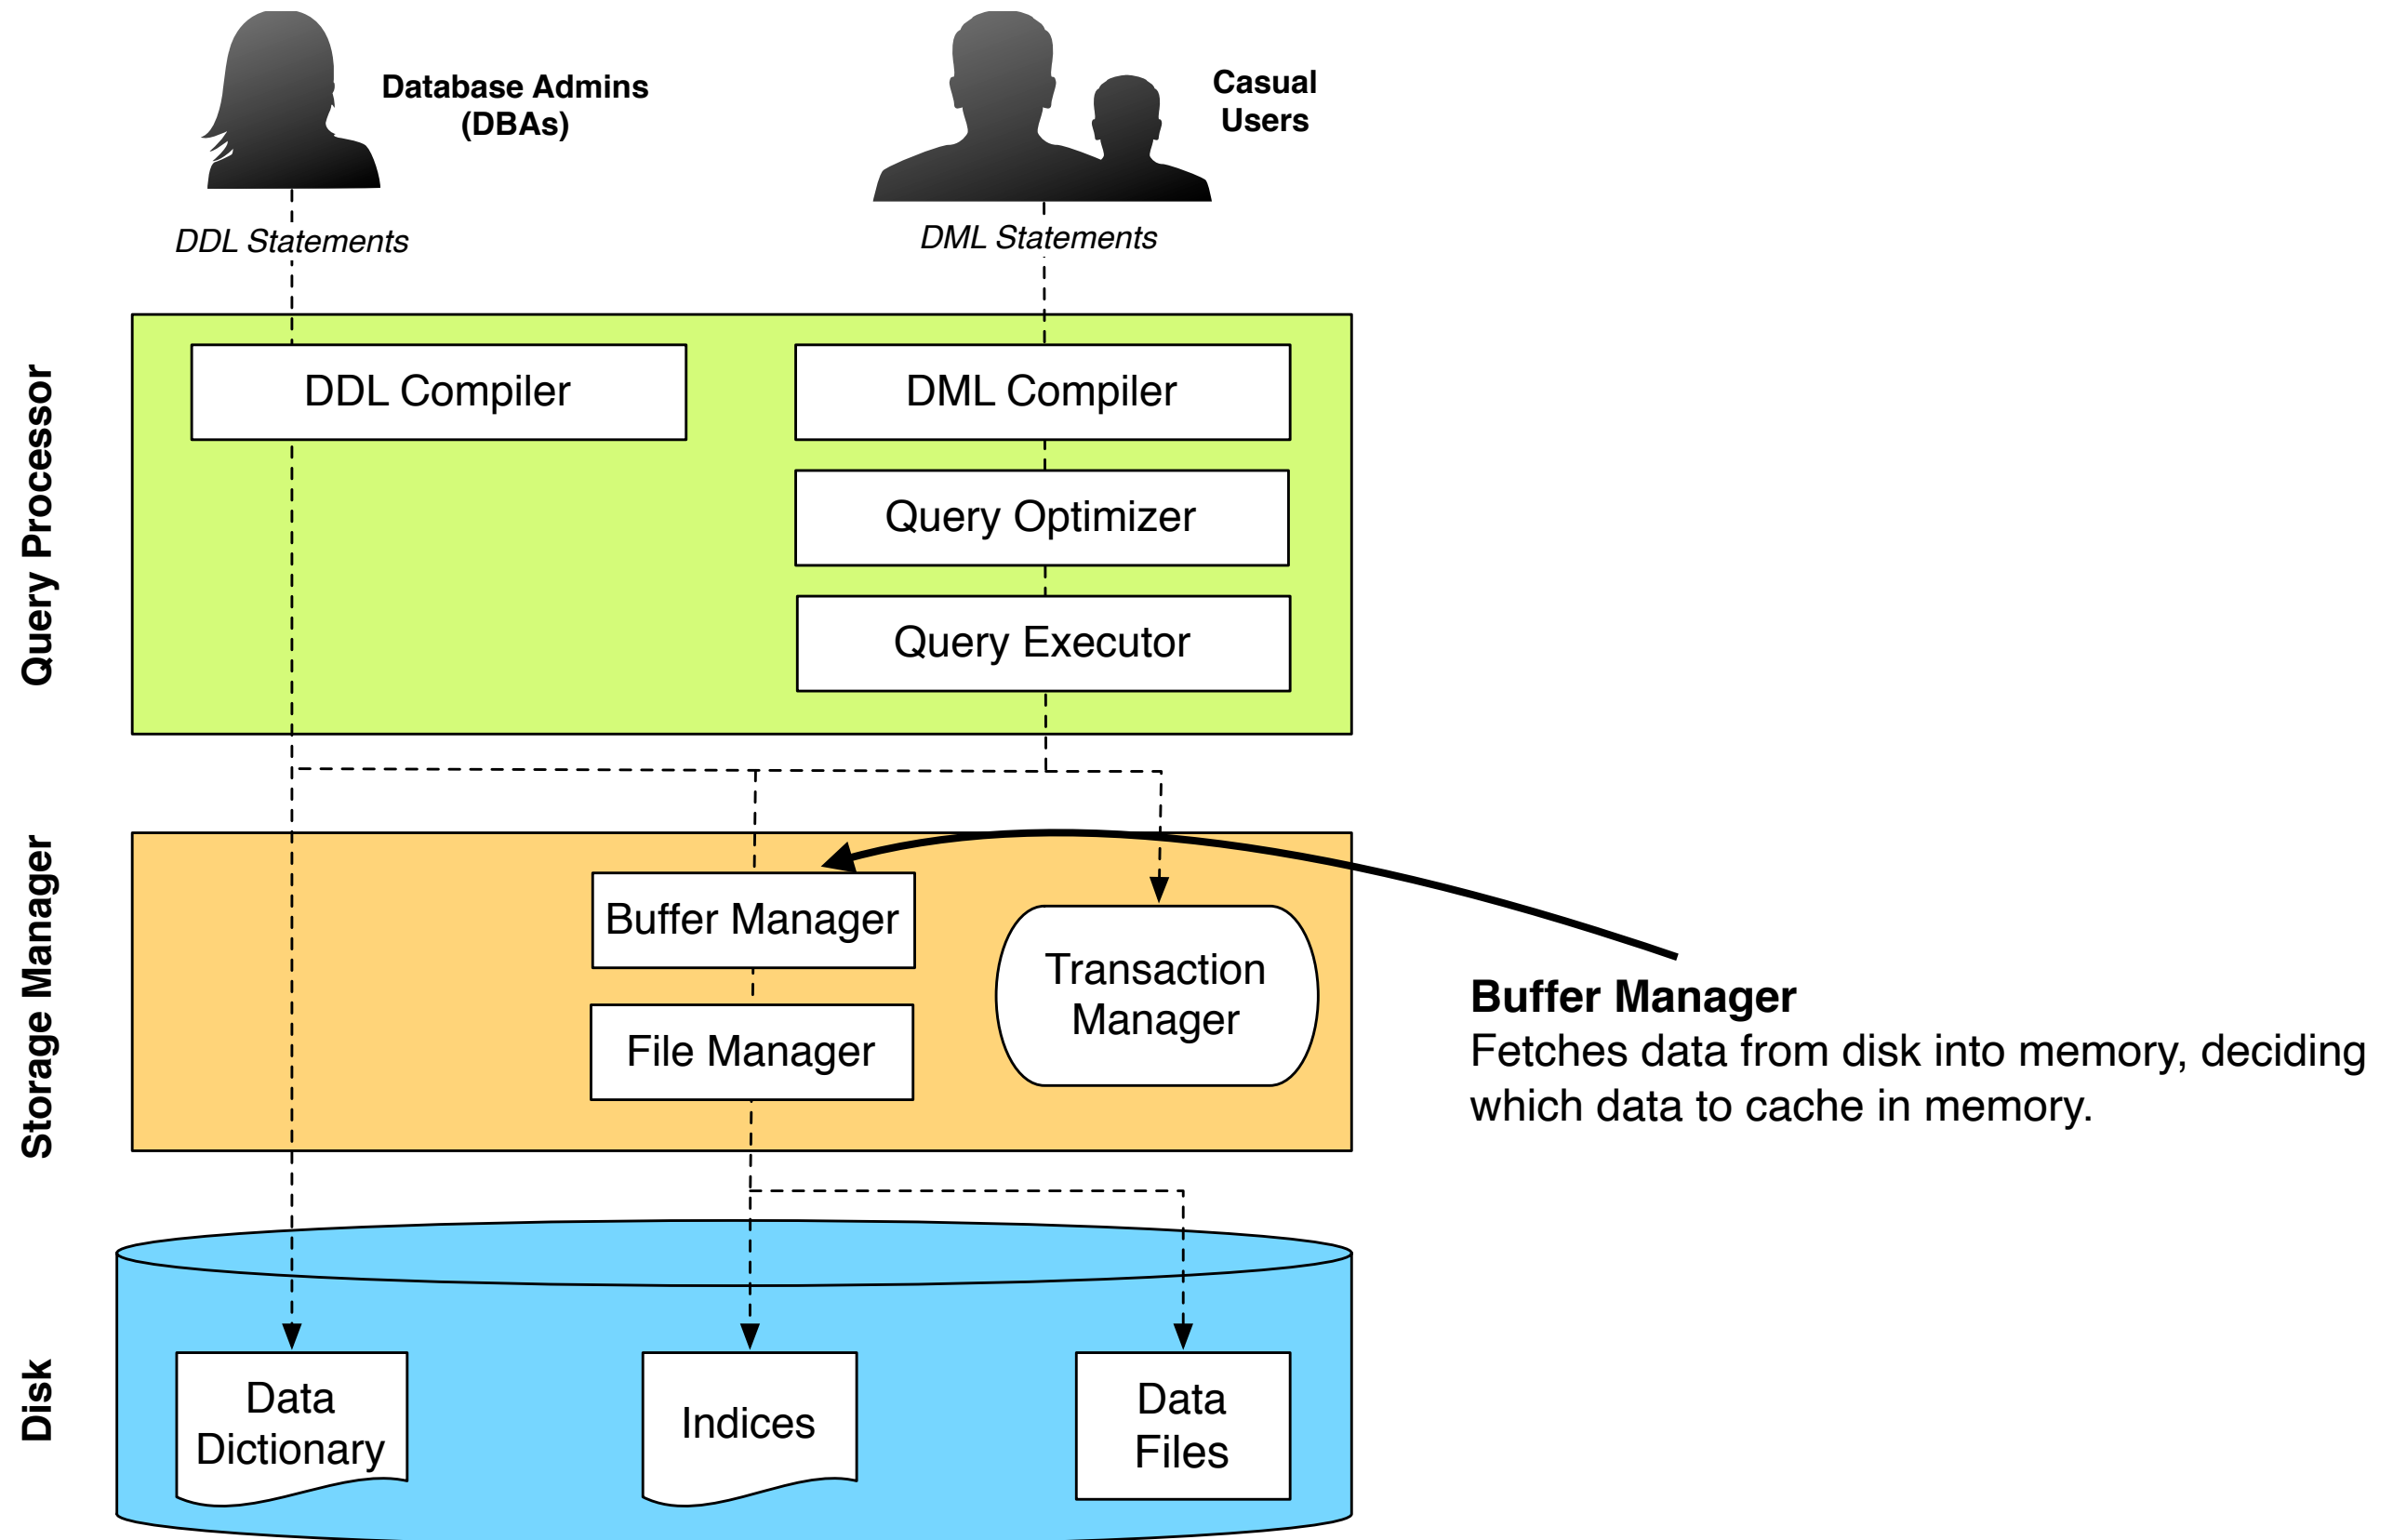
► Disks

- Hardware
- Understanding Files and Blocks
- Modeling Access Time
- Disk Scheduling

► Buffer Management



Database Architecture



Motivation: Buffer Management

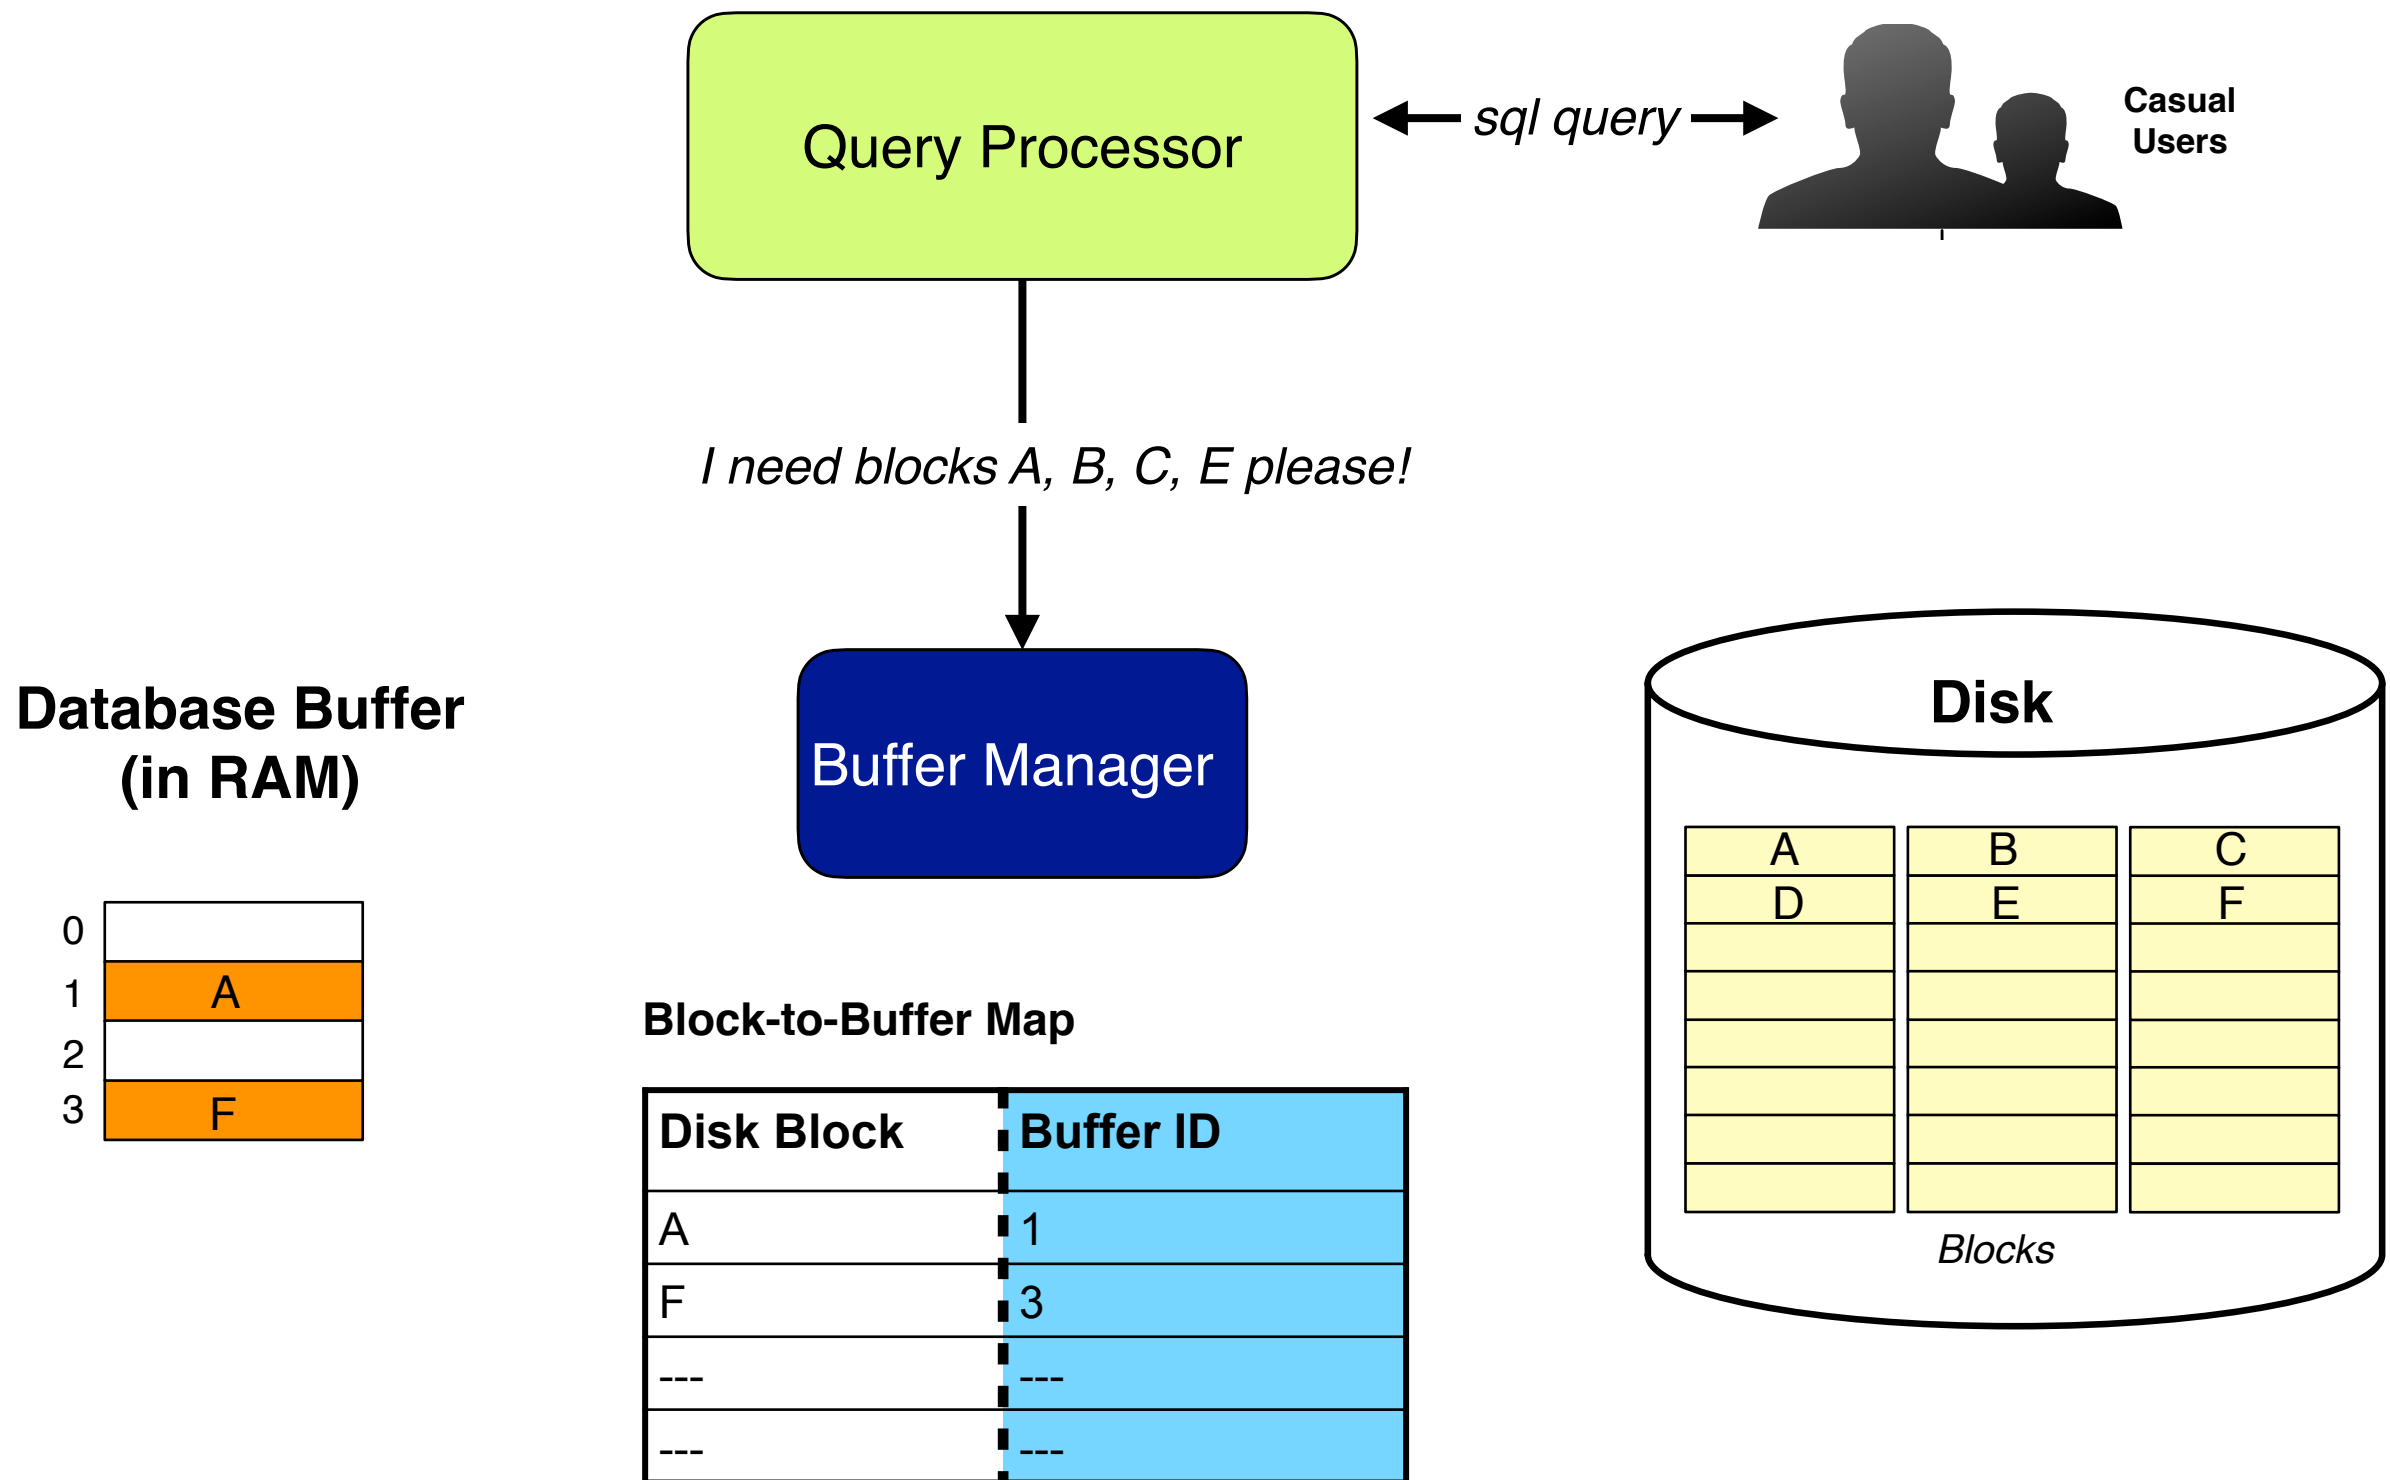
► Problem:

- Database size (on disk) is usually larger than available memory!
 - Terabytes (Disk) compared to GBs (RAM)
 - ~3 orders of magnitude difference!

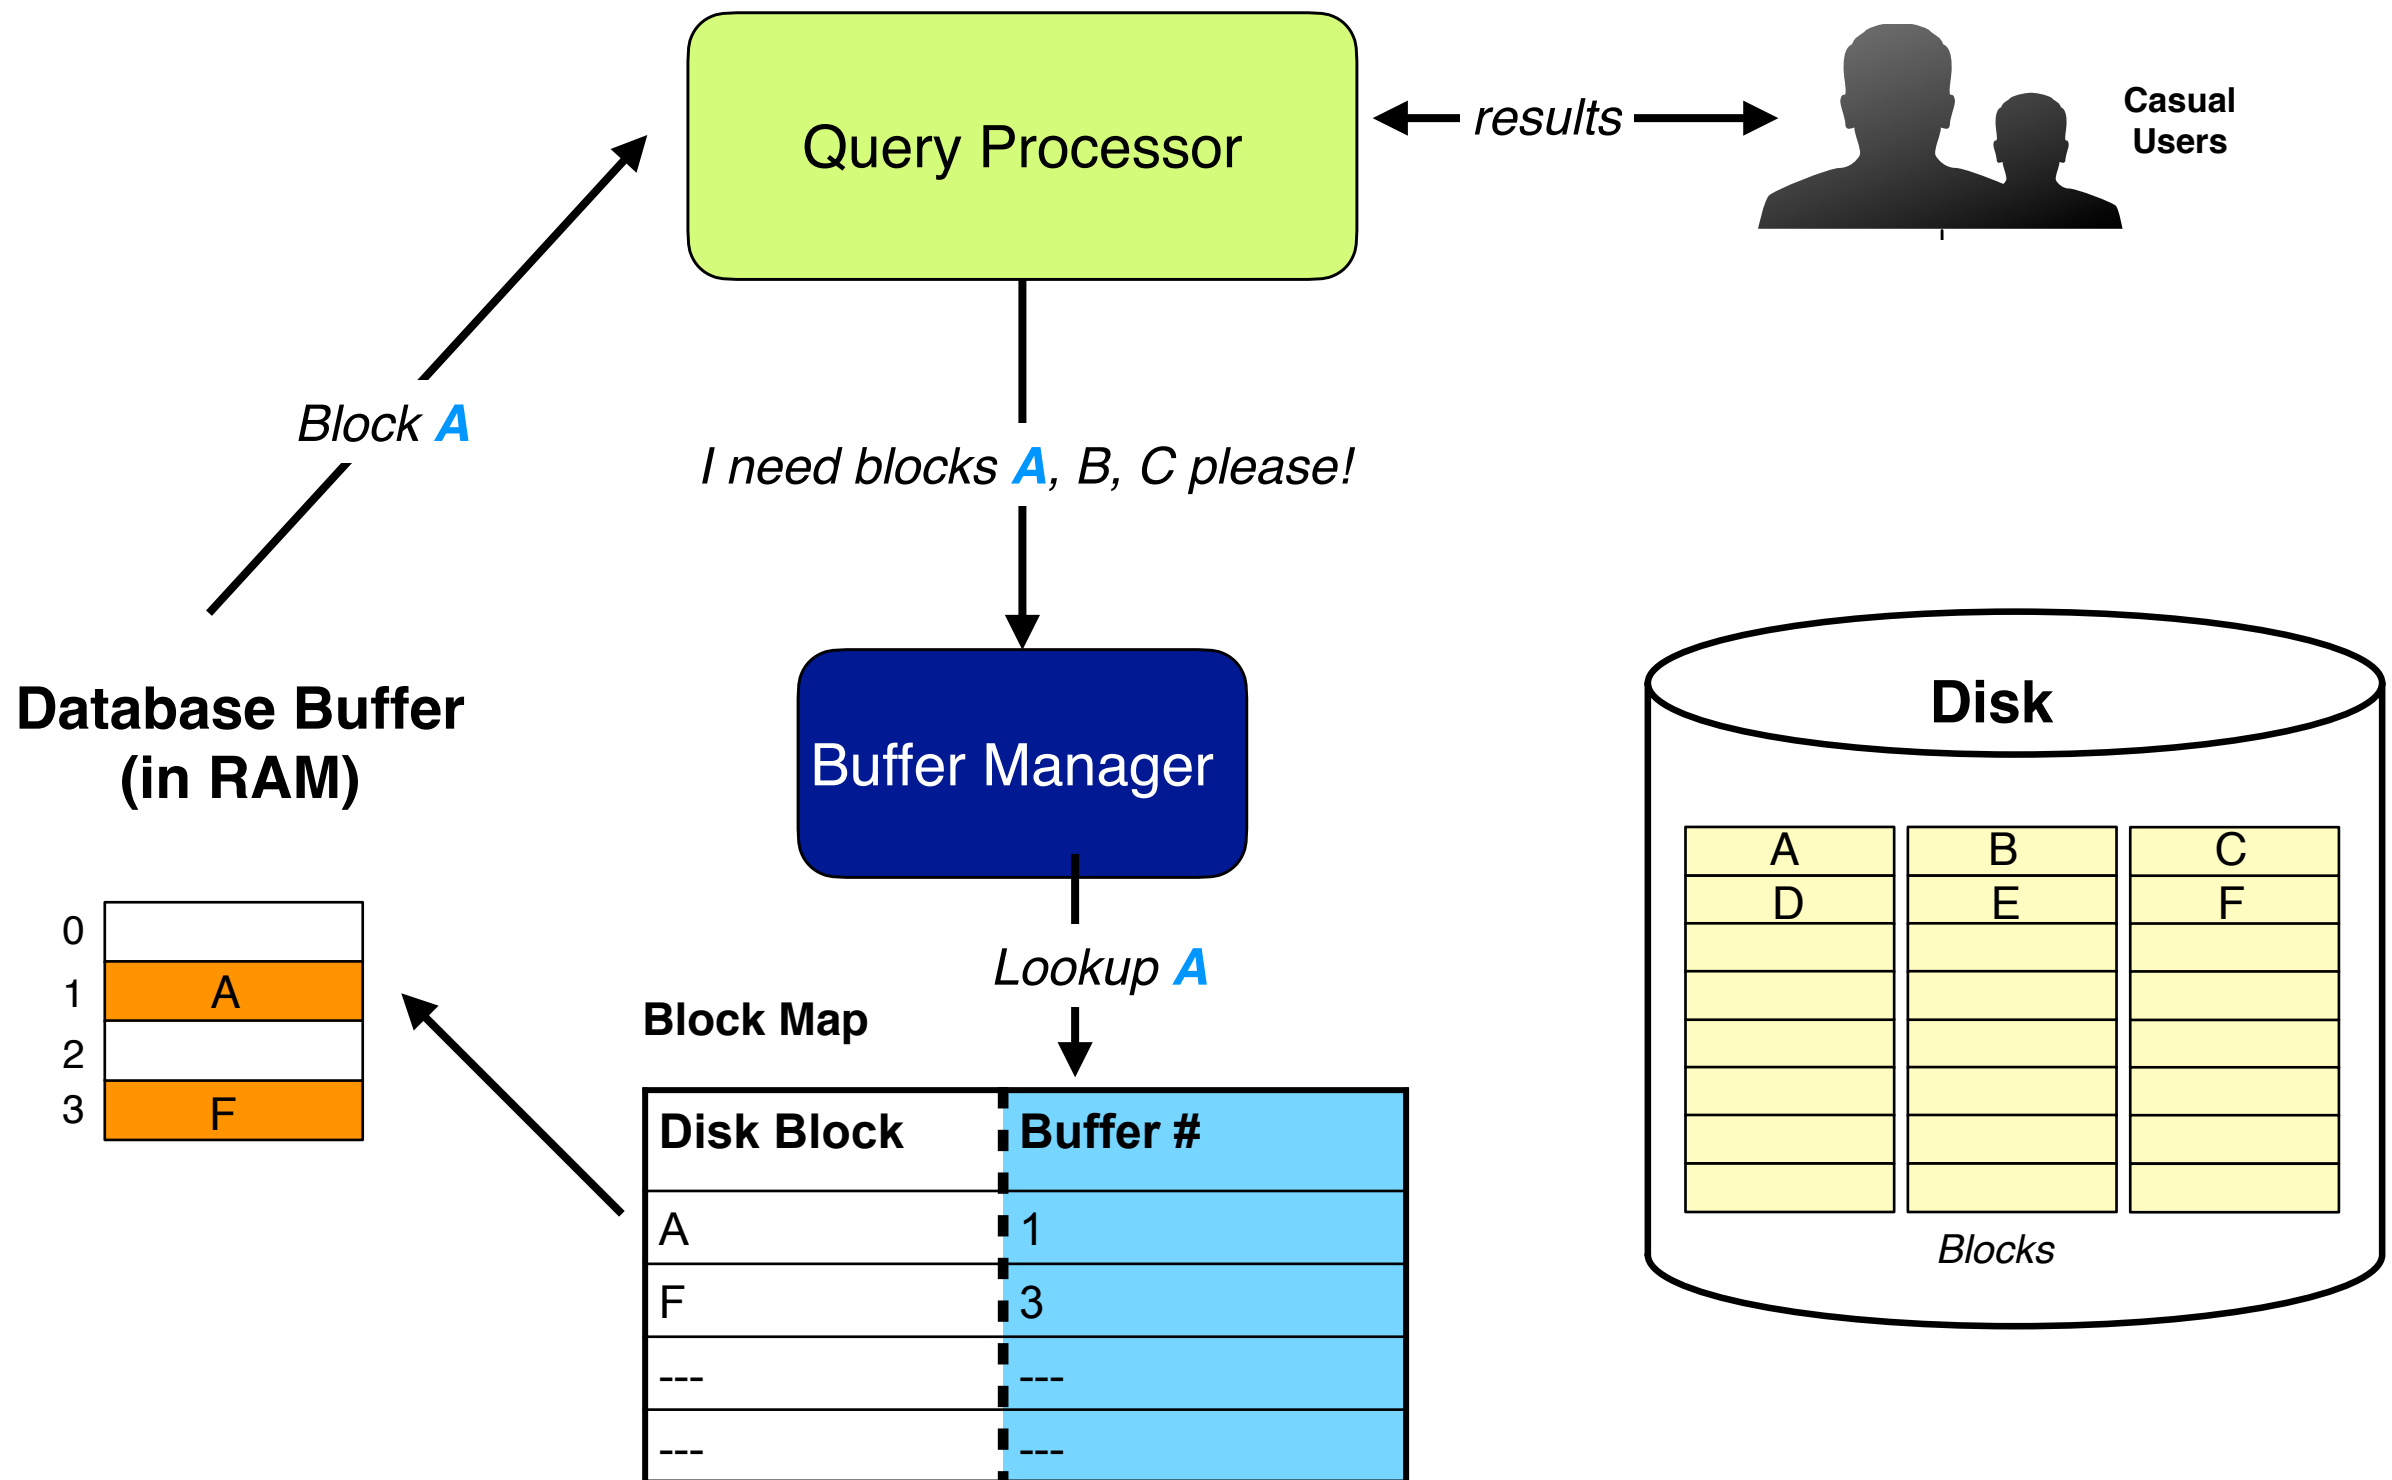
► Goal: Minimize number of block transfers between disk and memory while processing queries.

► A *database buffer* is a segment of memory that the DB uses to cache some disk blocks for later access.

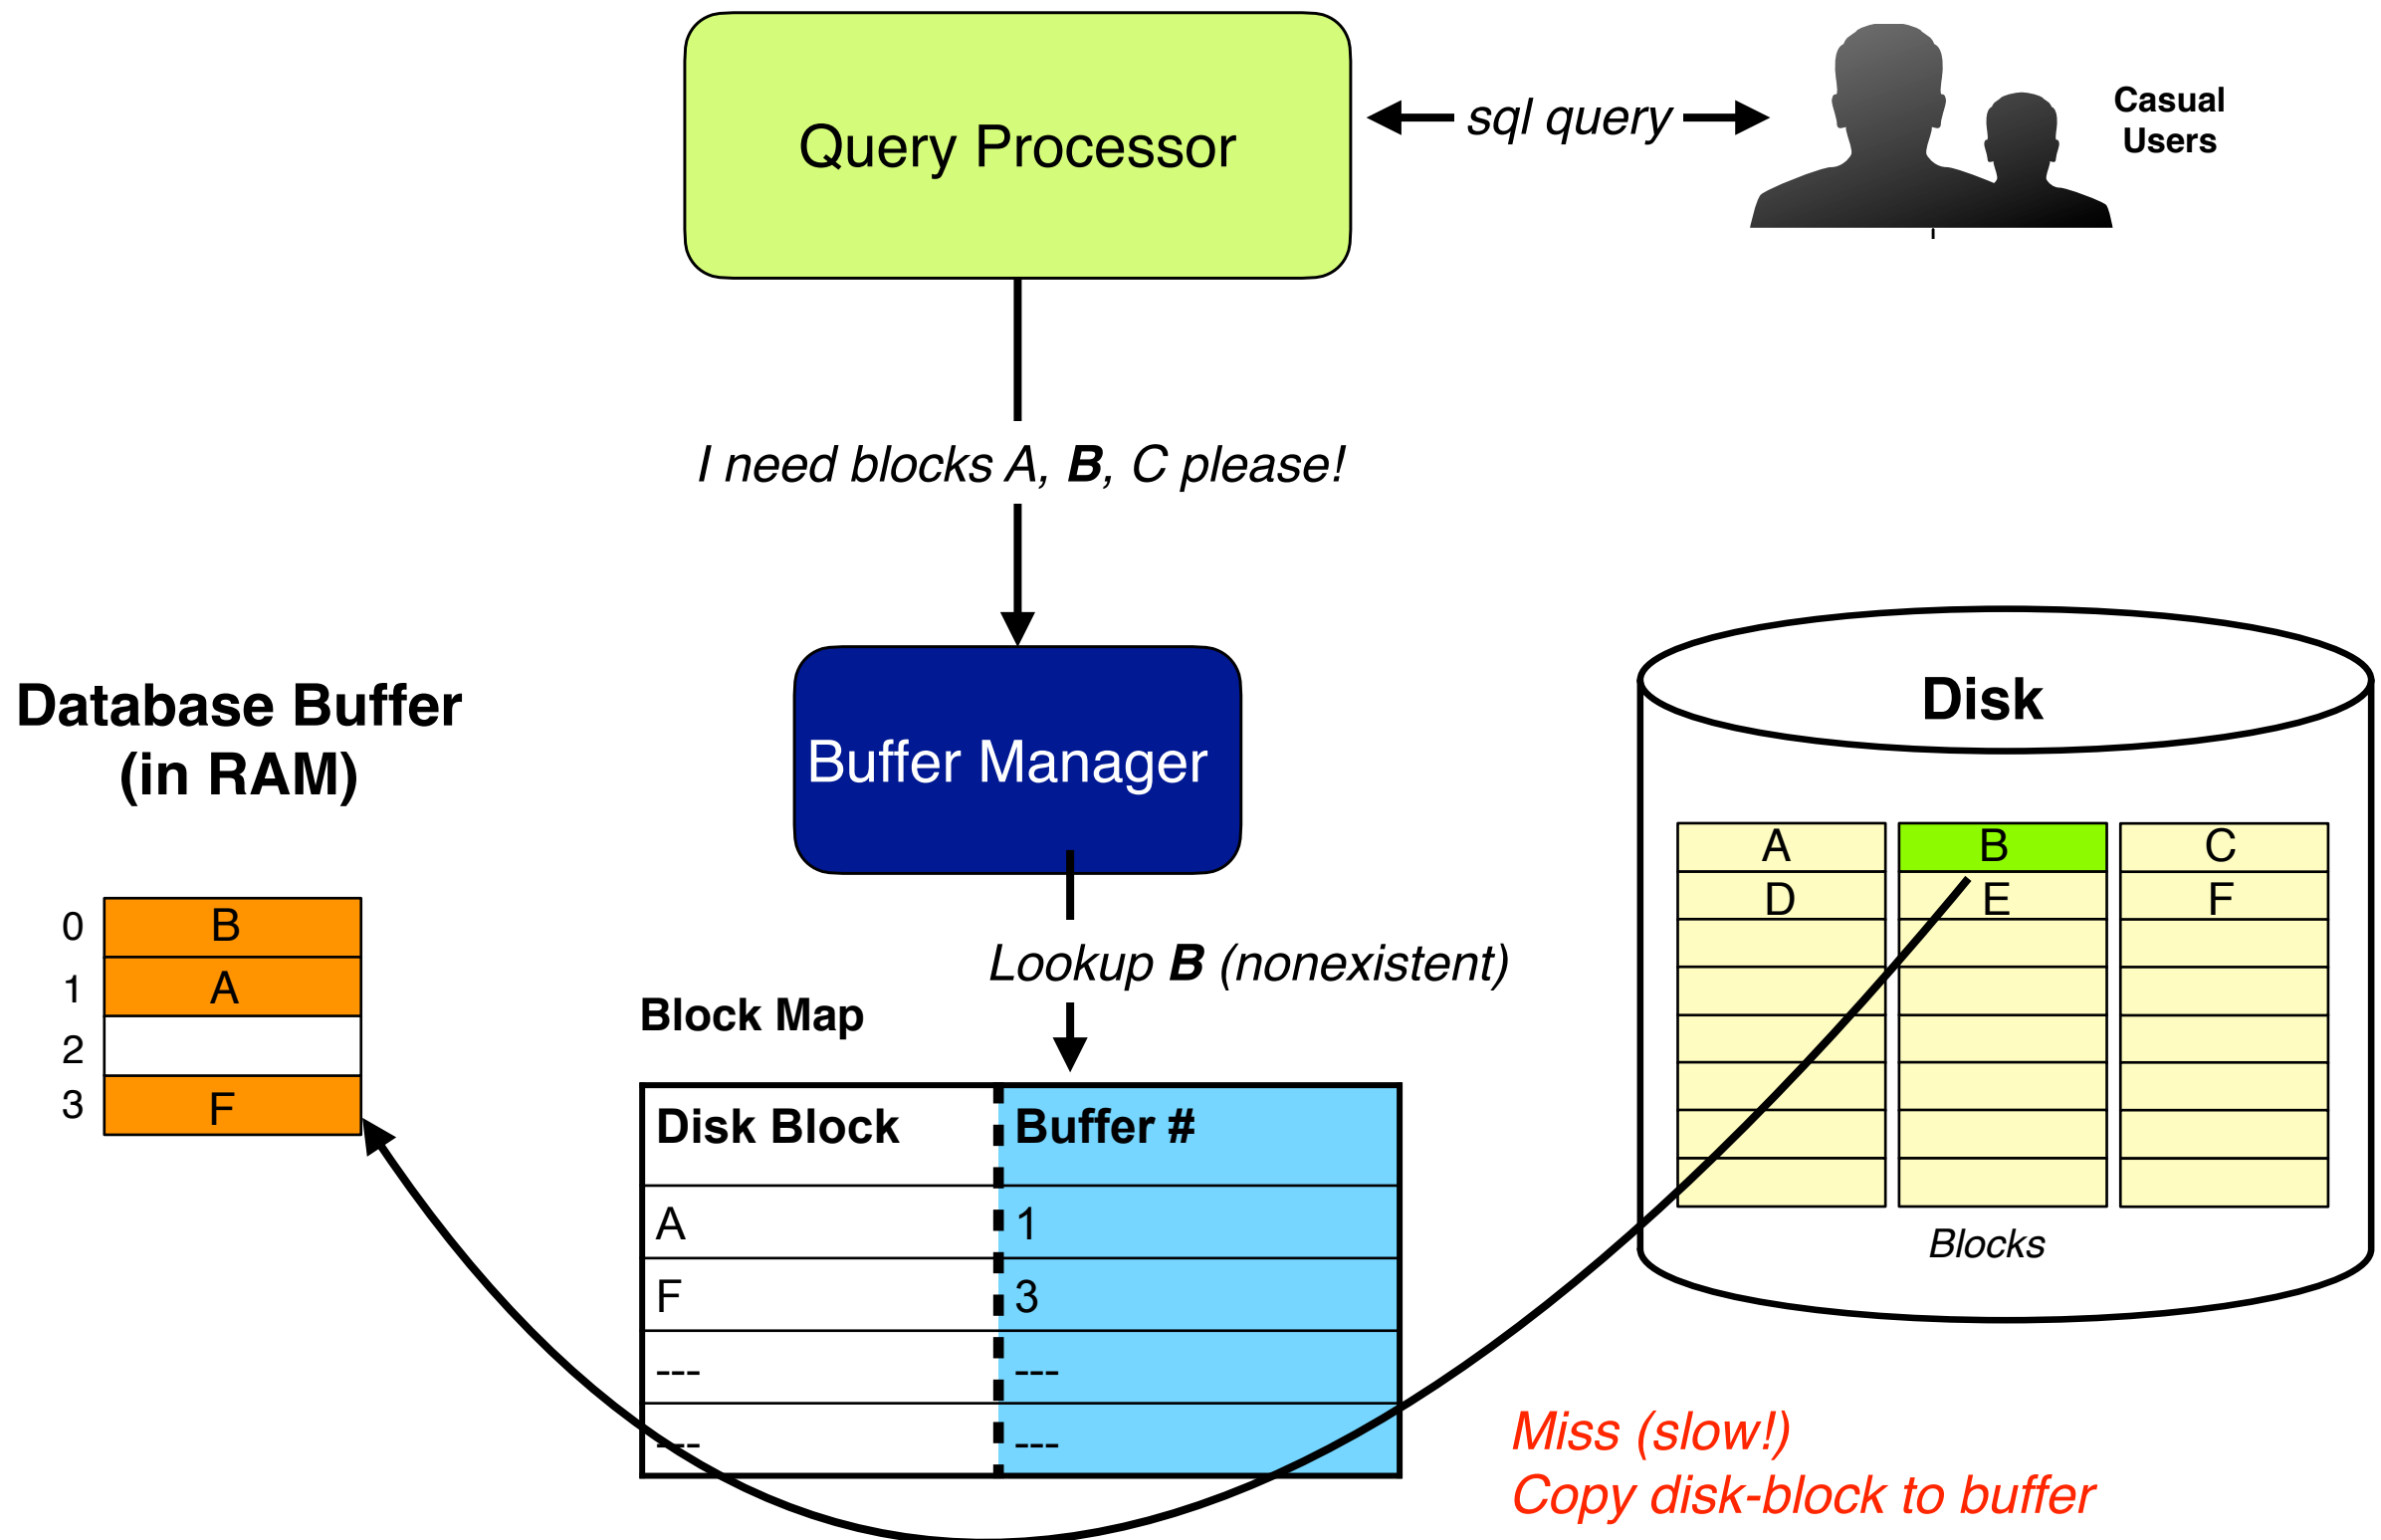
Buffer Manager



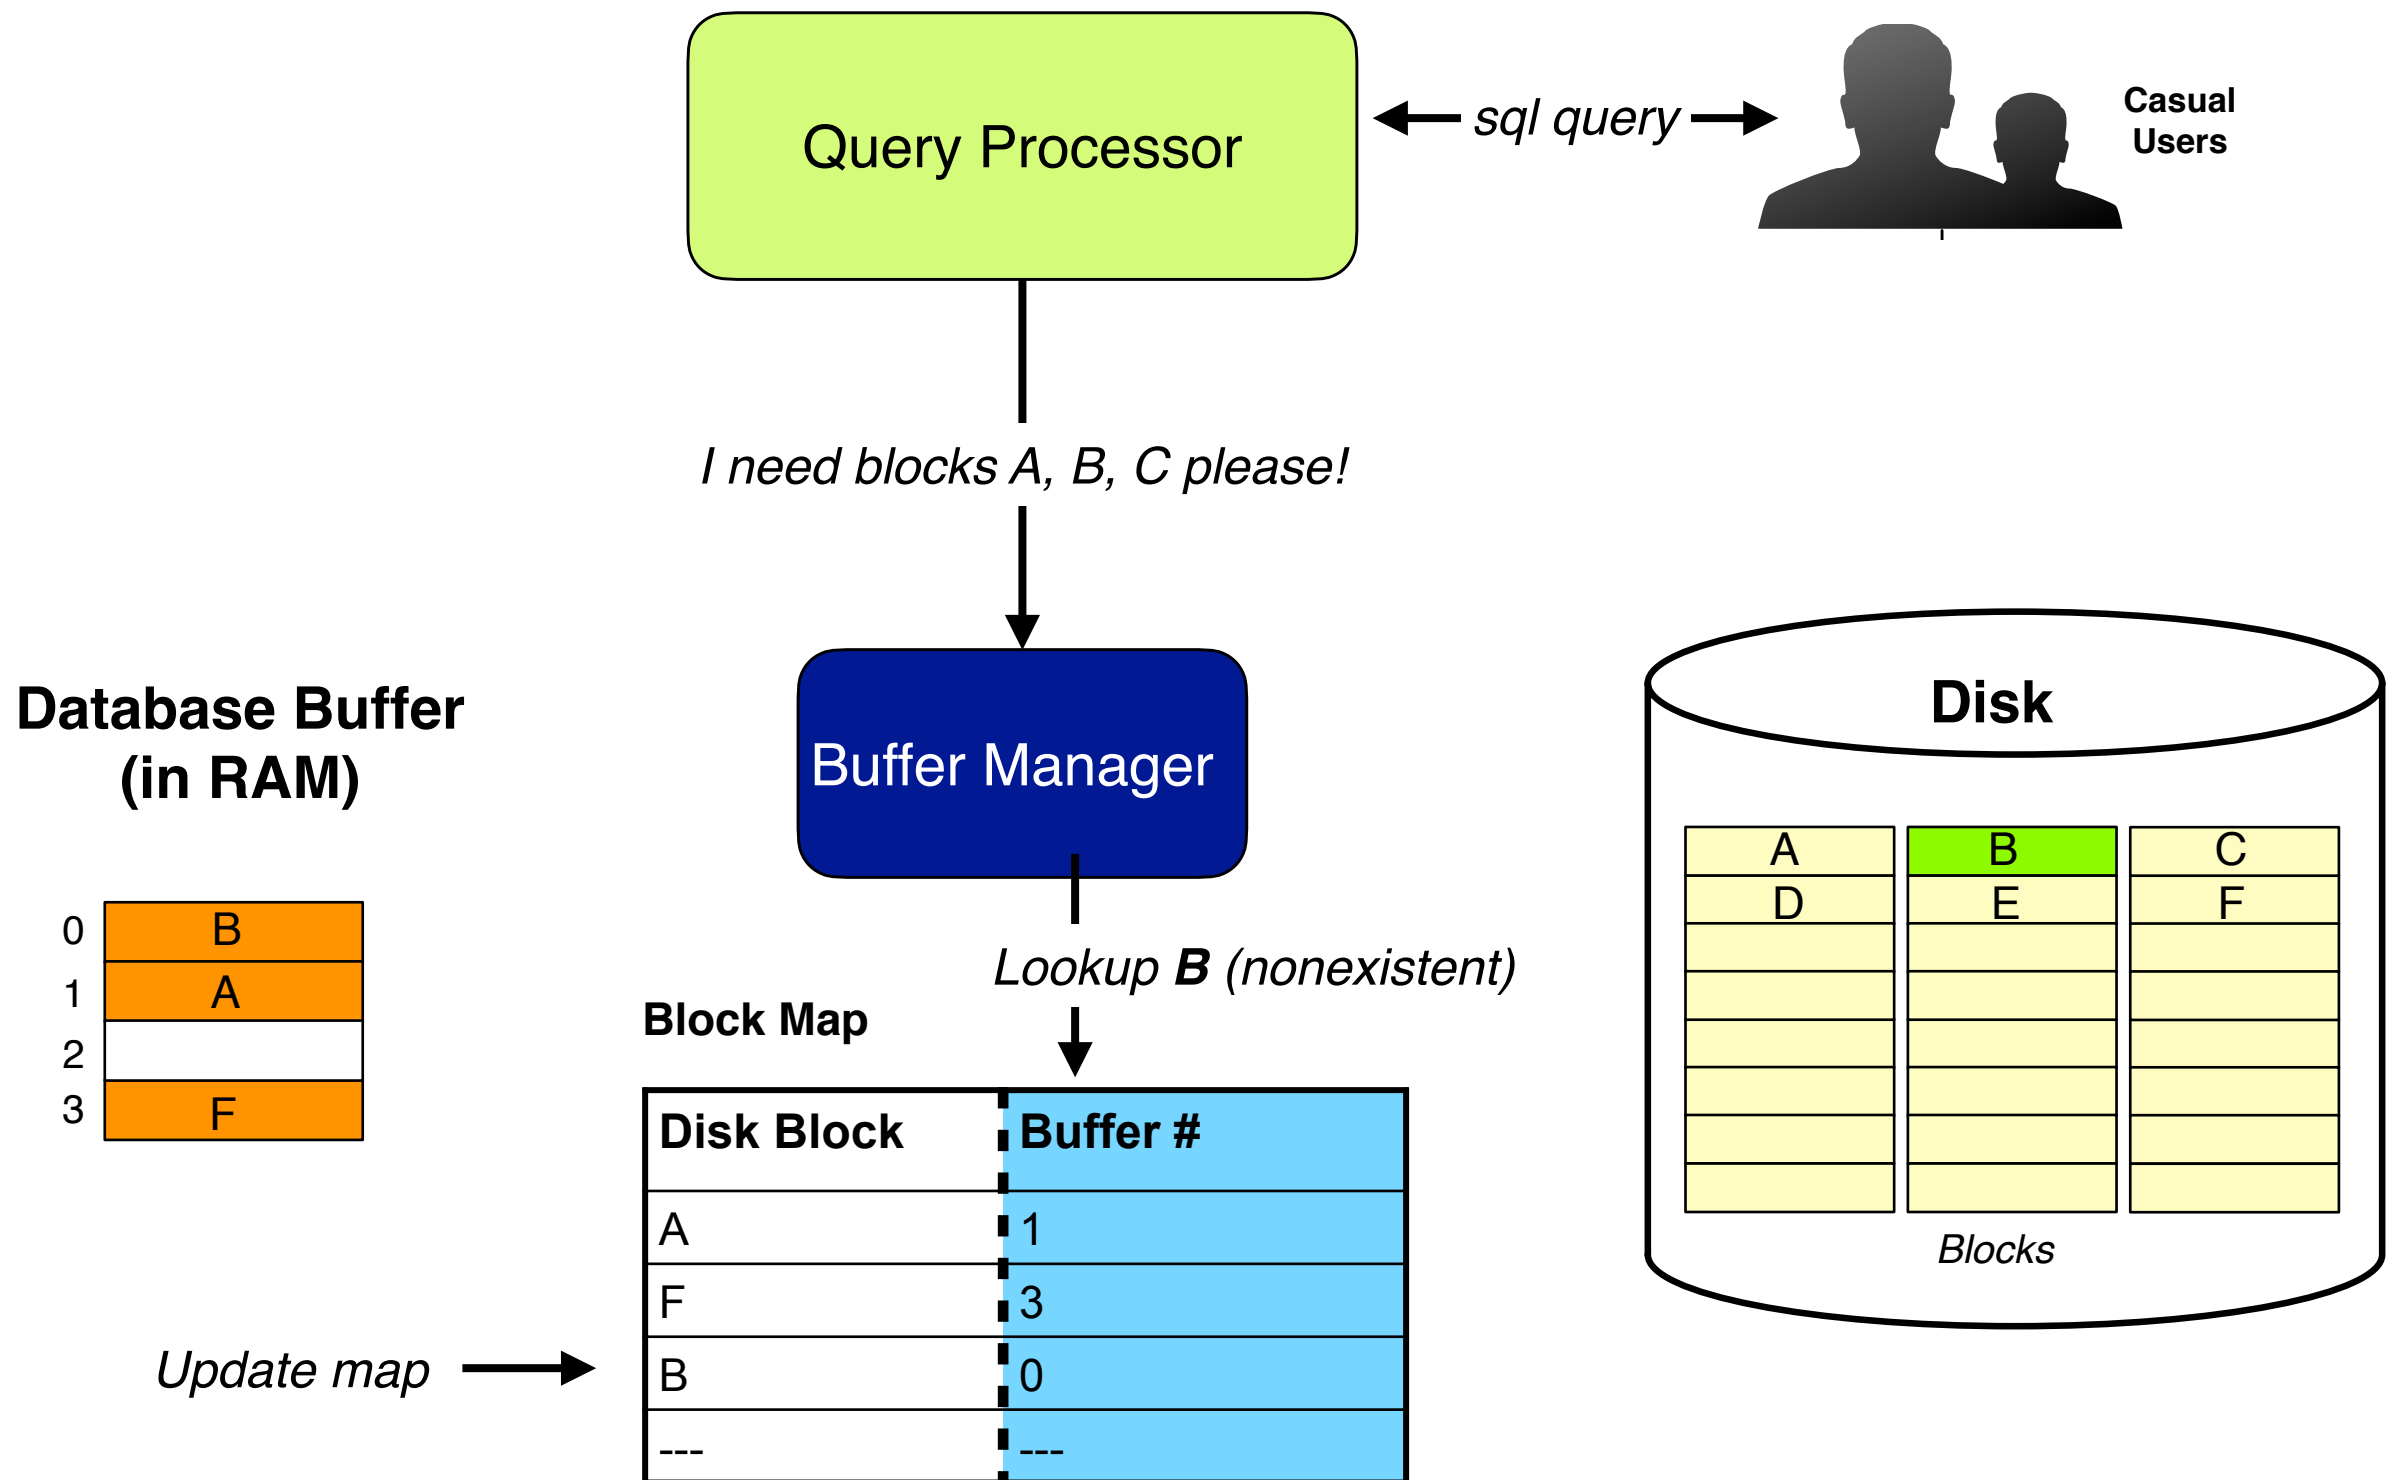
Buffer Manager



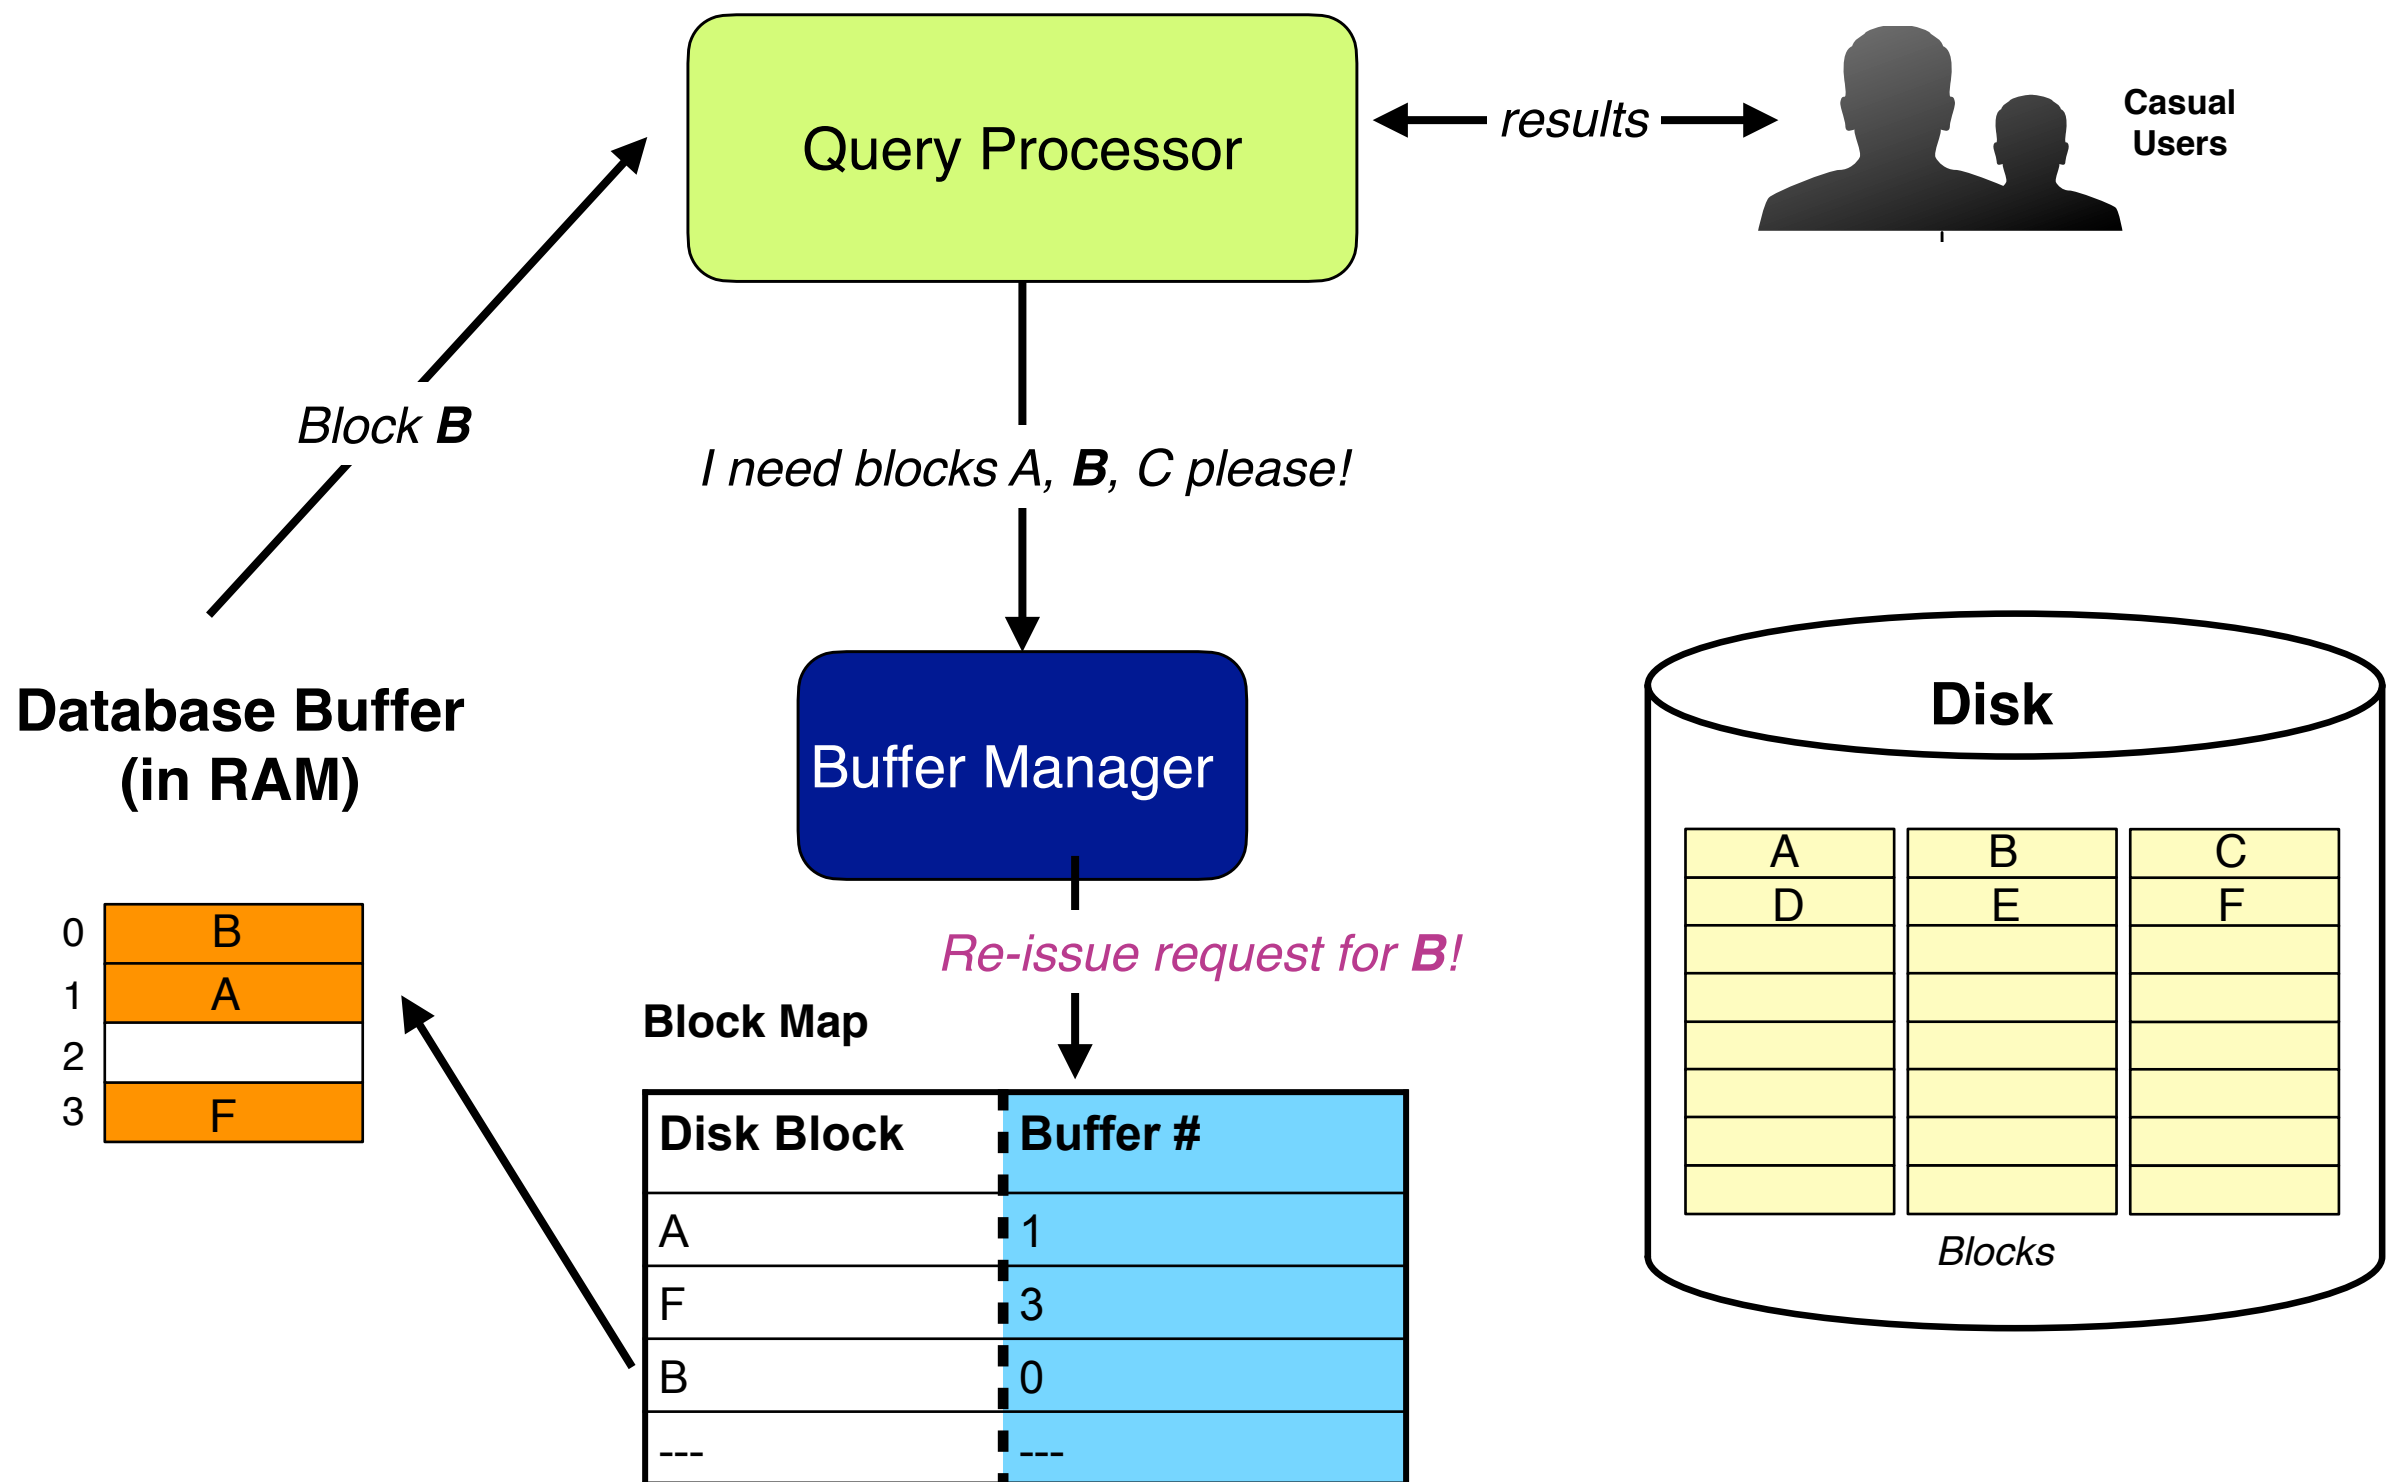
Buffer Manager



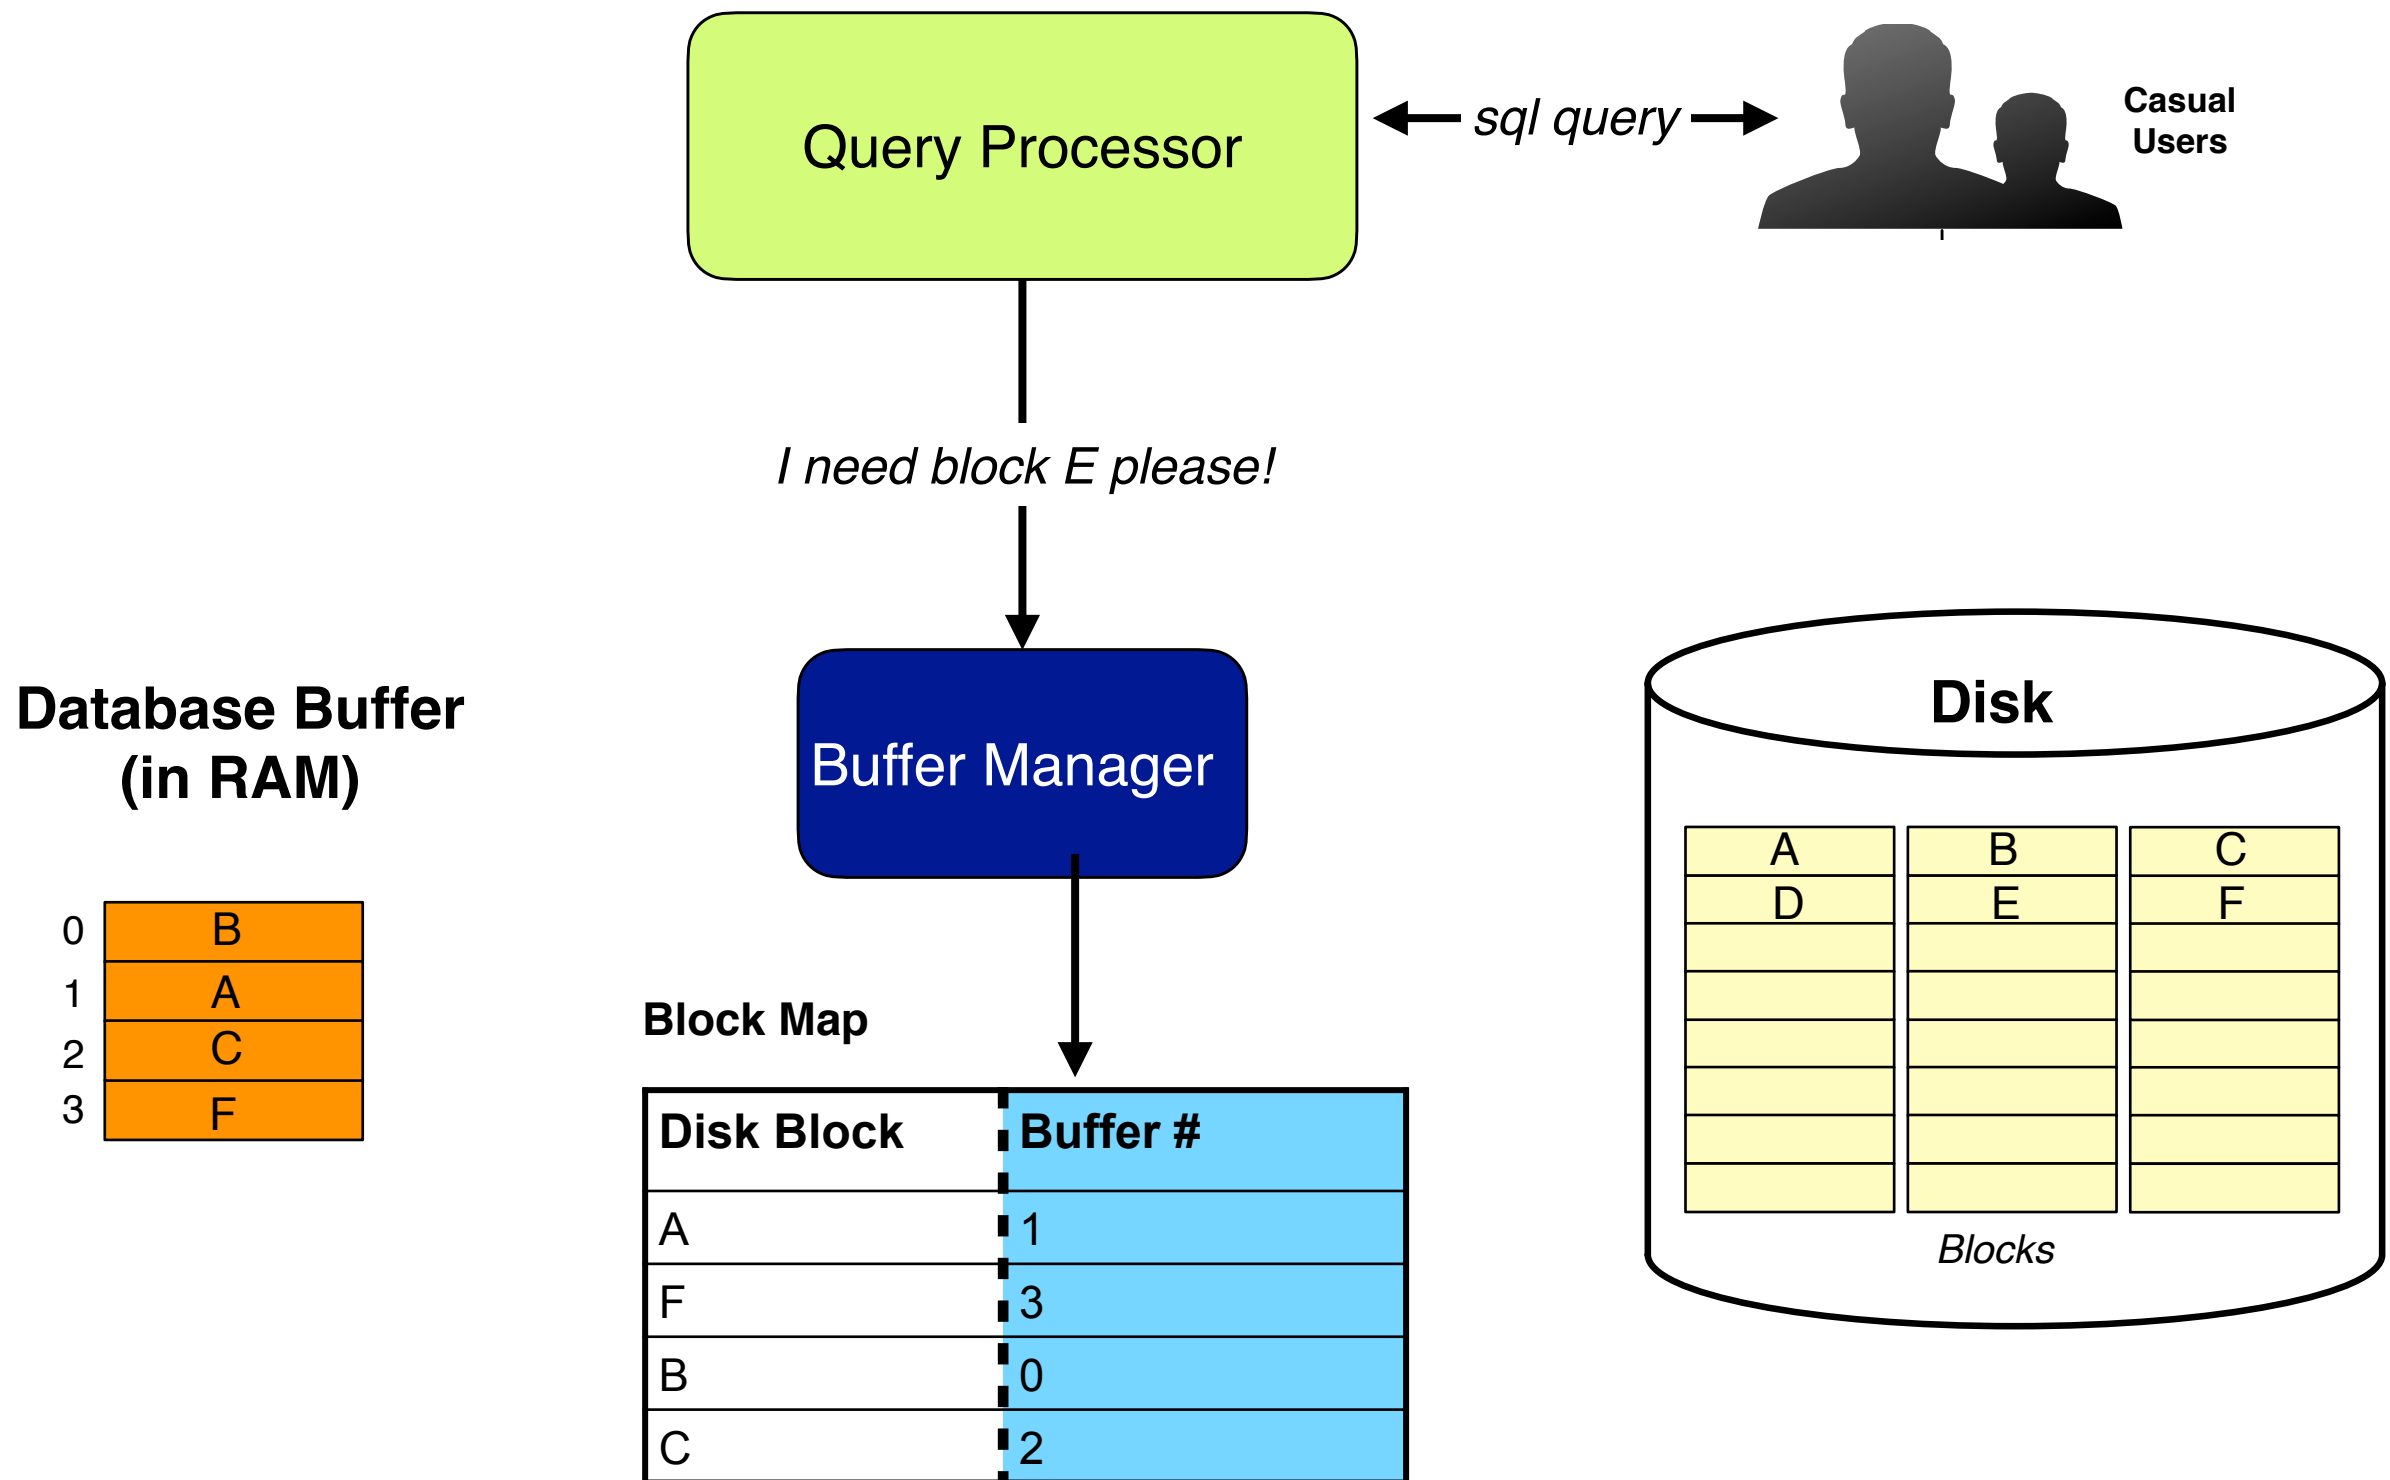
Buffer Manager



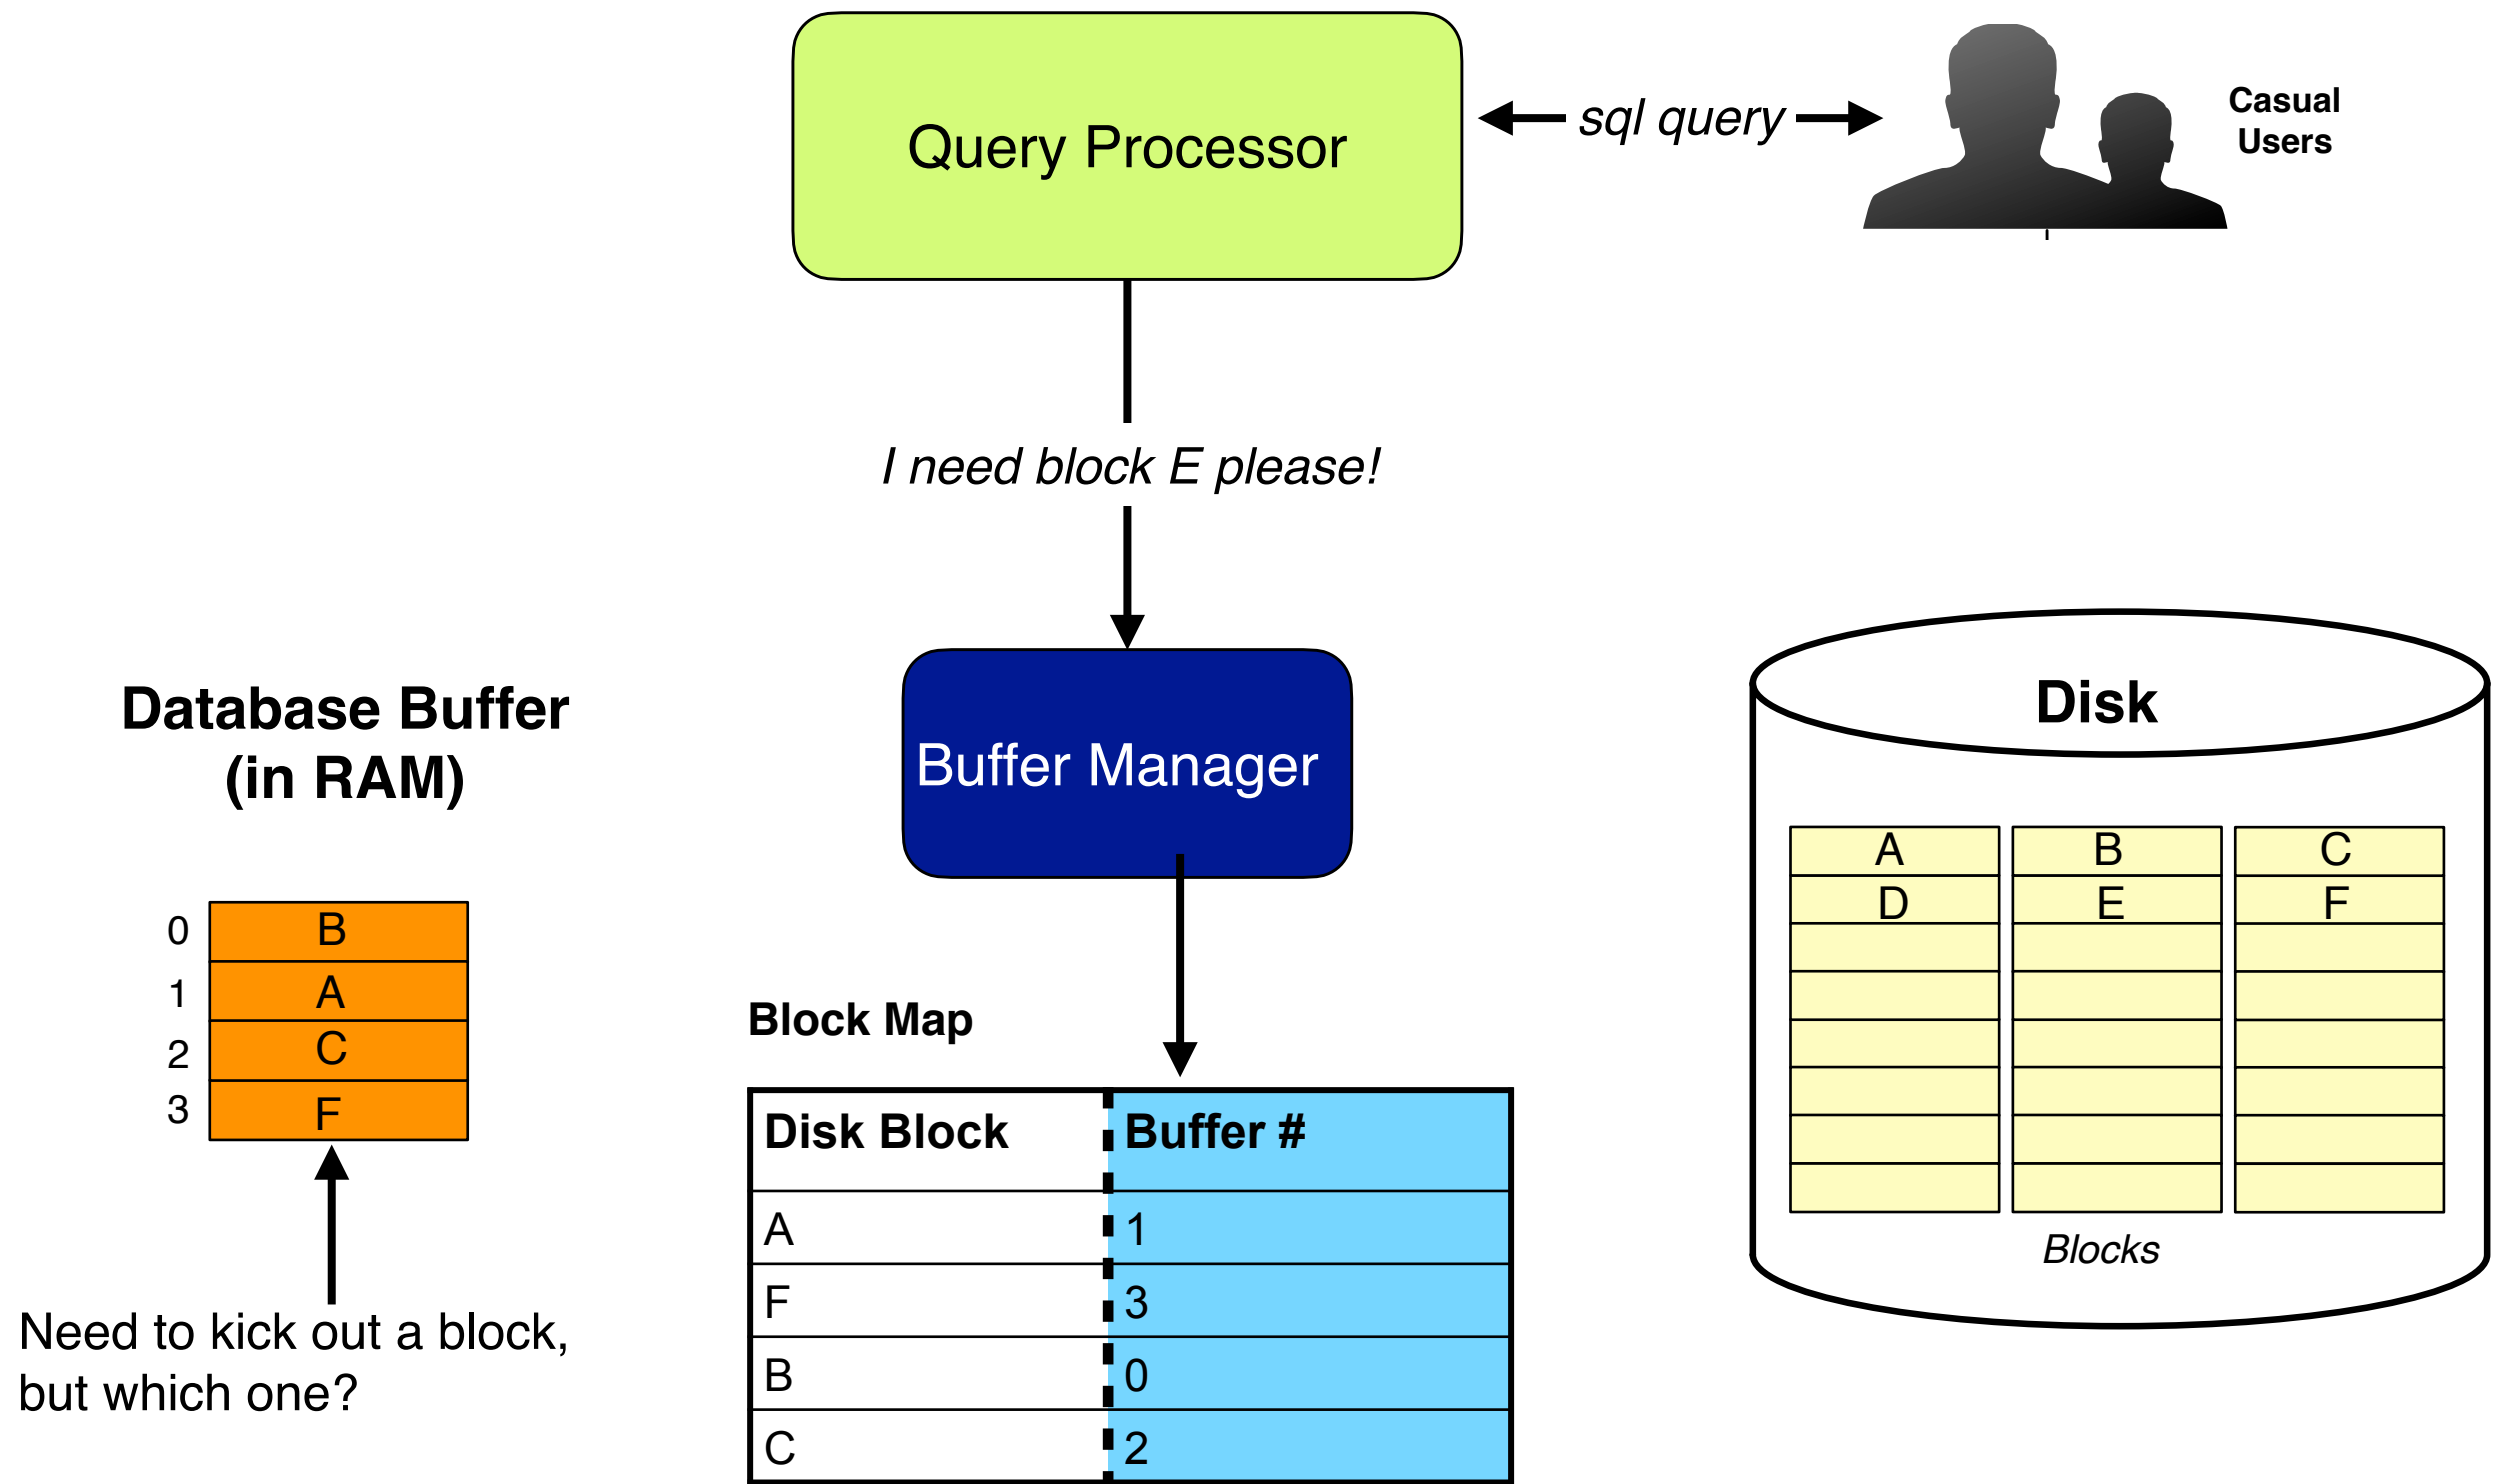
Buffer Manager



New Problem: What If Buffer Is Full?



New Problem: What If Buffer Is Full?



(Replacement Policy)

Common Block-Replacement Policies

- ▶ What does Buffer Manager do when buffer is full?
- ▶ *First In First Out (FIFO) Policy*
 - The oldest block allocated in the buffer gets kicked out.
 - Pros: Fast; Simple to implement using queue
 - Cons: An old block doesn't mean it's hardly-ever used!
- ▶ *Least Recently Used (LRU) Policy*
 - The block that was used farthest in the past gets kicked out.
 - Pros: Exploits temporal locality
 - Cons: Harder to implement (uses priority queue)

Example

- ▶ Assume we only have *2 blocks* in the database buffer.
- ▶ Two tables: **Book**(BNO, title, author) and **Publisher**(PNO, BNO, Name)
 - 100 Books tuples can be stored in a block
 - 50 Publisher tuples can be stored in a block

Book	

	BNO title author
Block B1	book1
	...
	book100
Block B2	book101
	...
	book200
Block B3	book201
	...
	book300

Publisher	

	PNO BNO Name
Block P1	pub1
	...
	pub50
Block P2	pub51
	...
	pub100

Example (Cont.)

► Let's run a simple select query: `select * from Books;`


► Access Pattern:

	Books

	BN0 title author
Block B1	book1
	...
	book100
Block B2	book101
	...
	book200
Block B3	book201
	...
	book300

FIFO Replacement Policy


- Database buffer allocates **two** blocks in RAM
 - 3 file blocks for the Books relation: B1, B2, B3
 - Access Pattern of *select_query()*: B1 (50x), B2 (50x), B3 (50x)



Number of misses to run query = 3

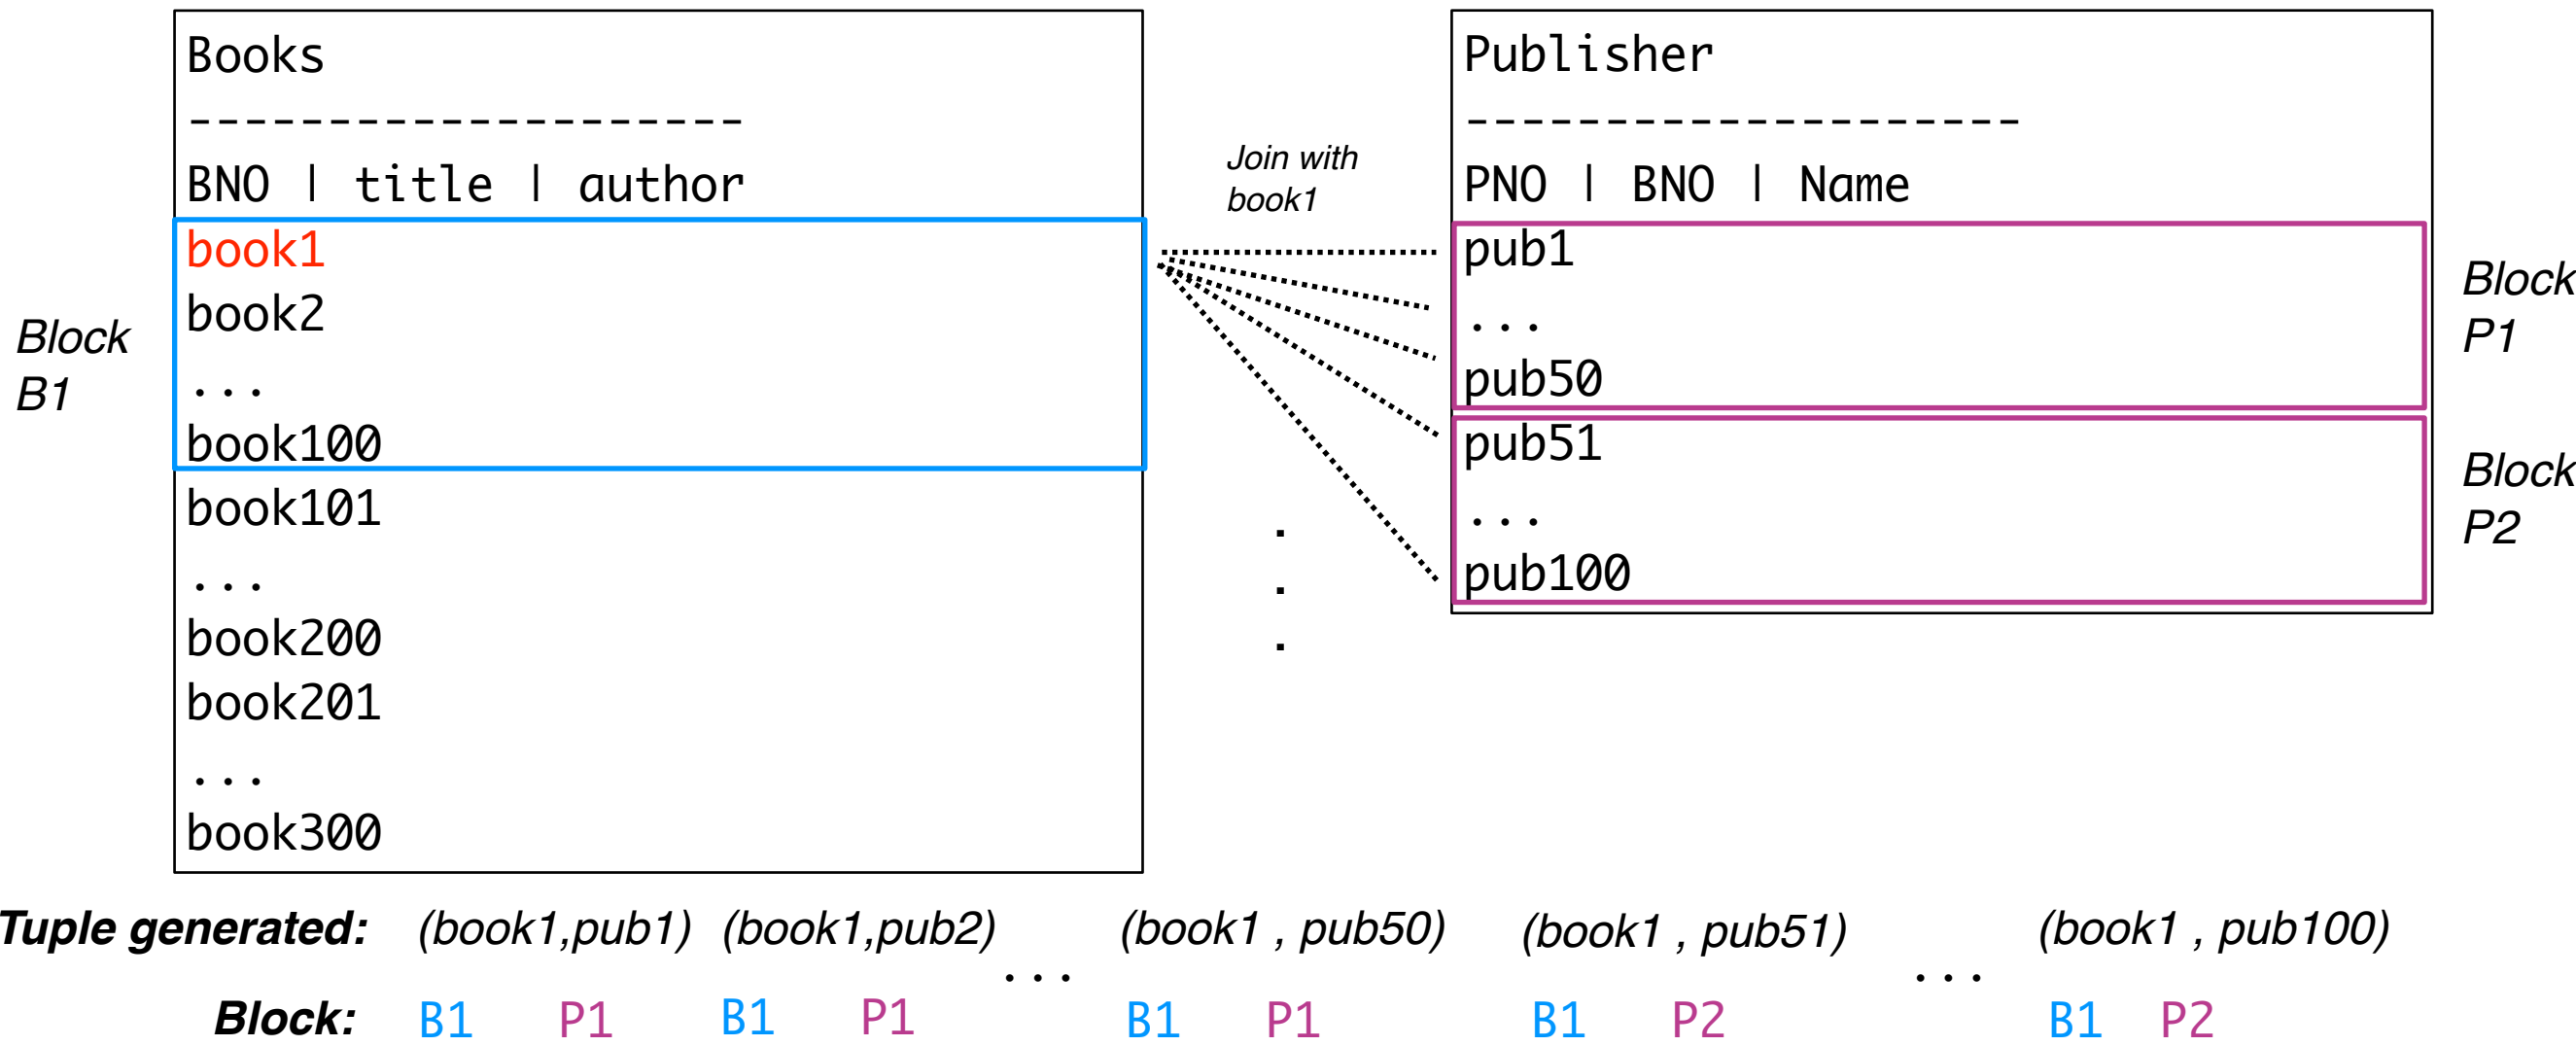
LRU Replacement

- Database buffer allocates **two** blocks in RAM
 - 3 file blocks for the Books relation: B1, B2, B3
 - Access Pattern of *select_query()*: B1 (50x), B2 (50x), B3 (50x)

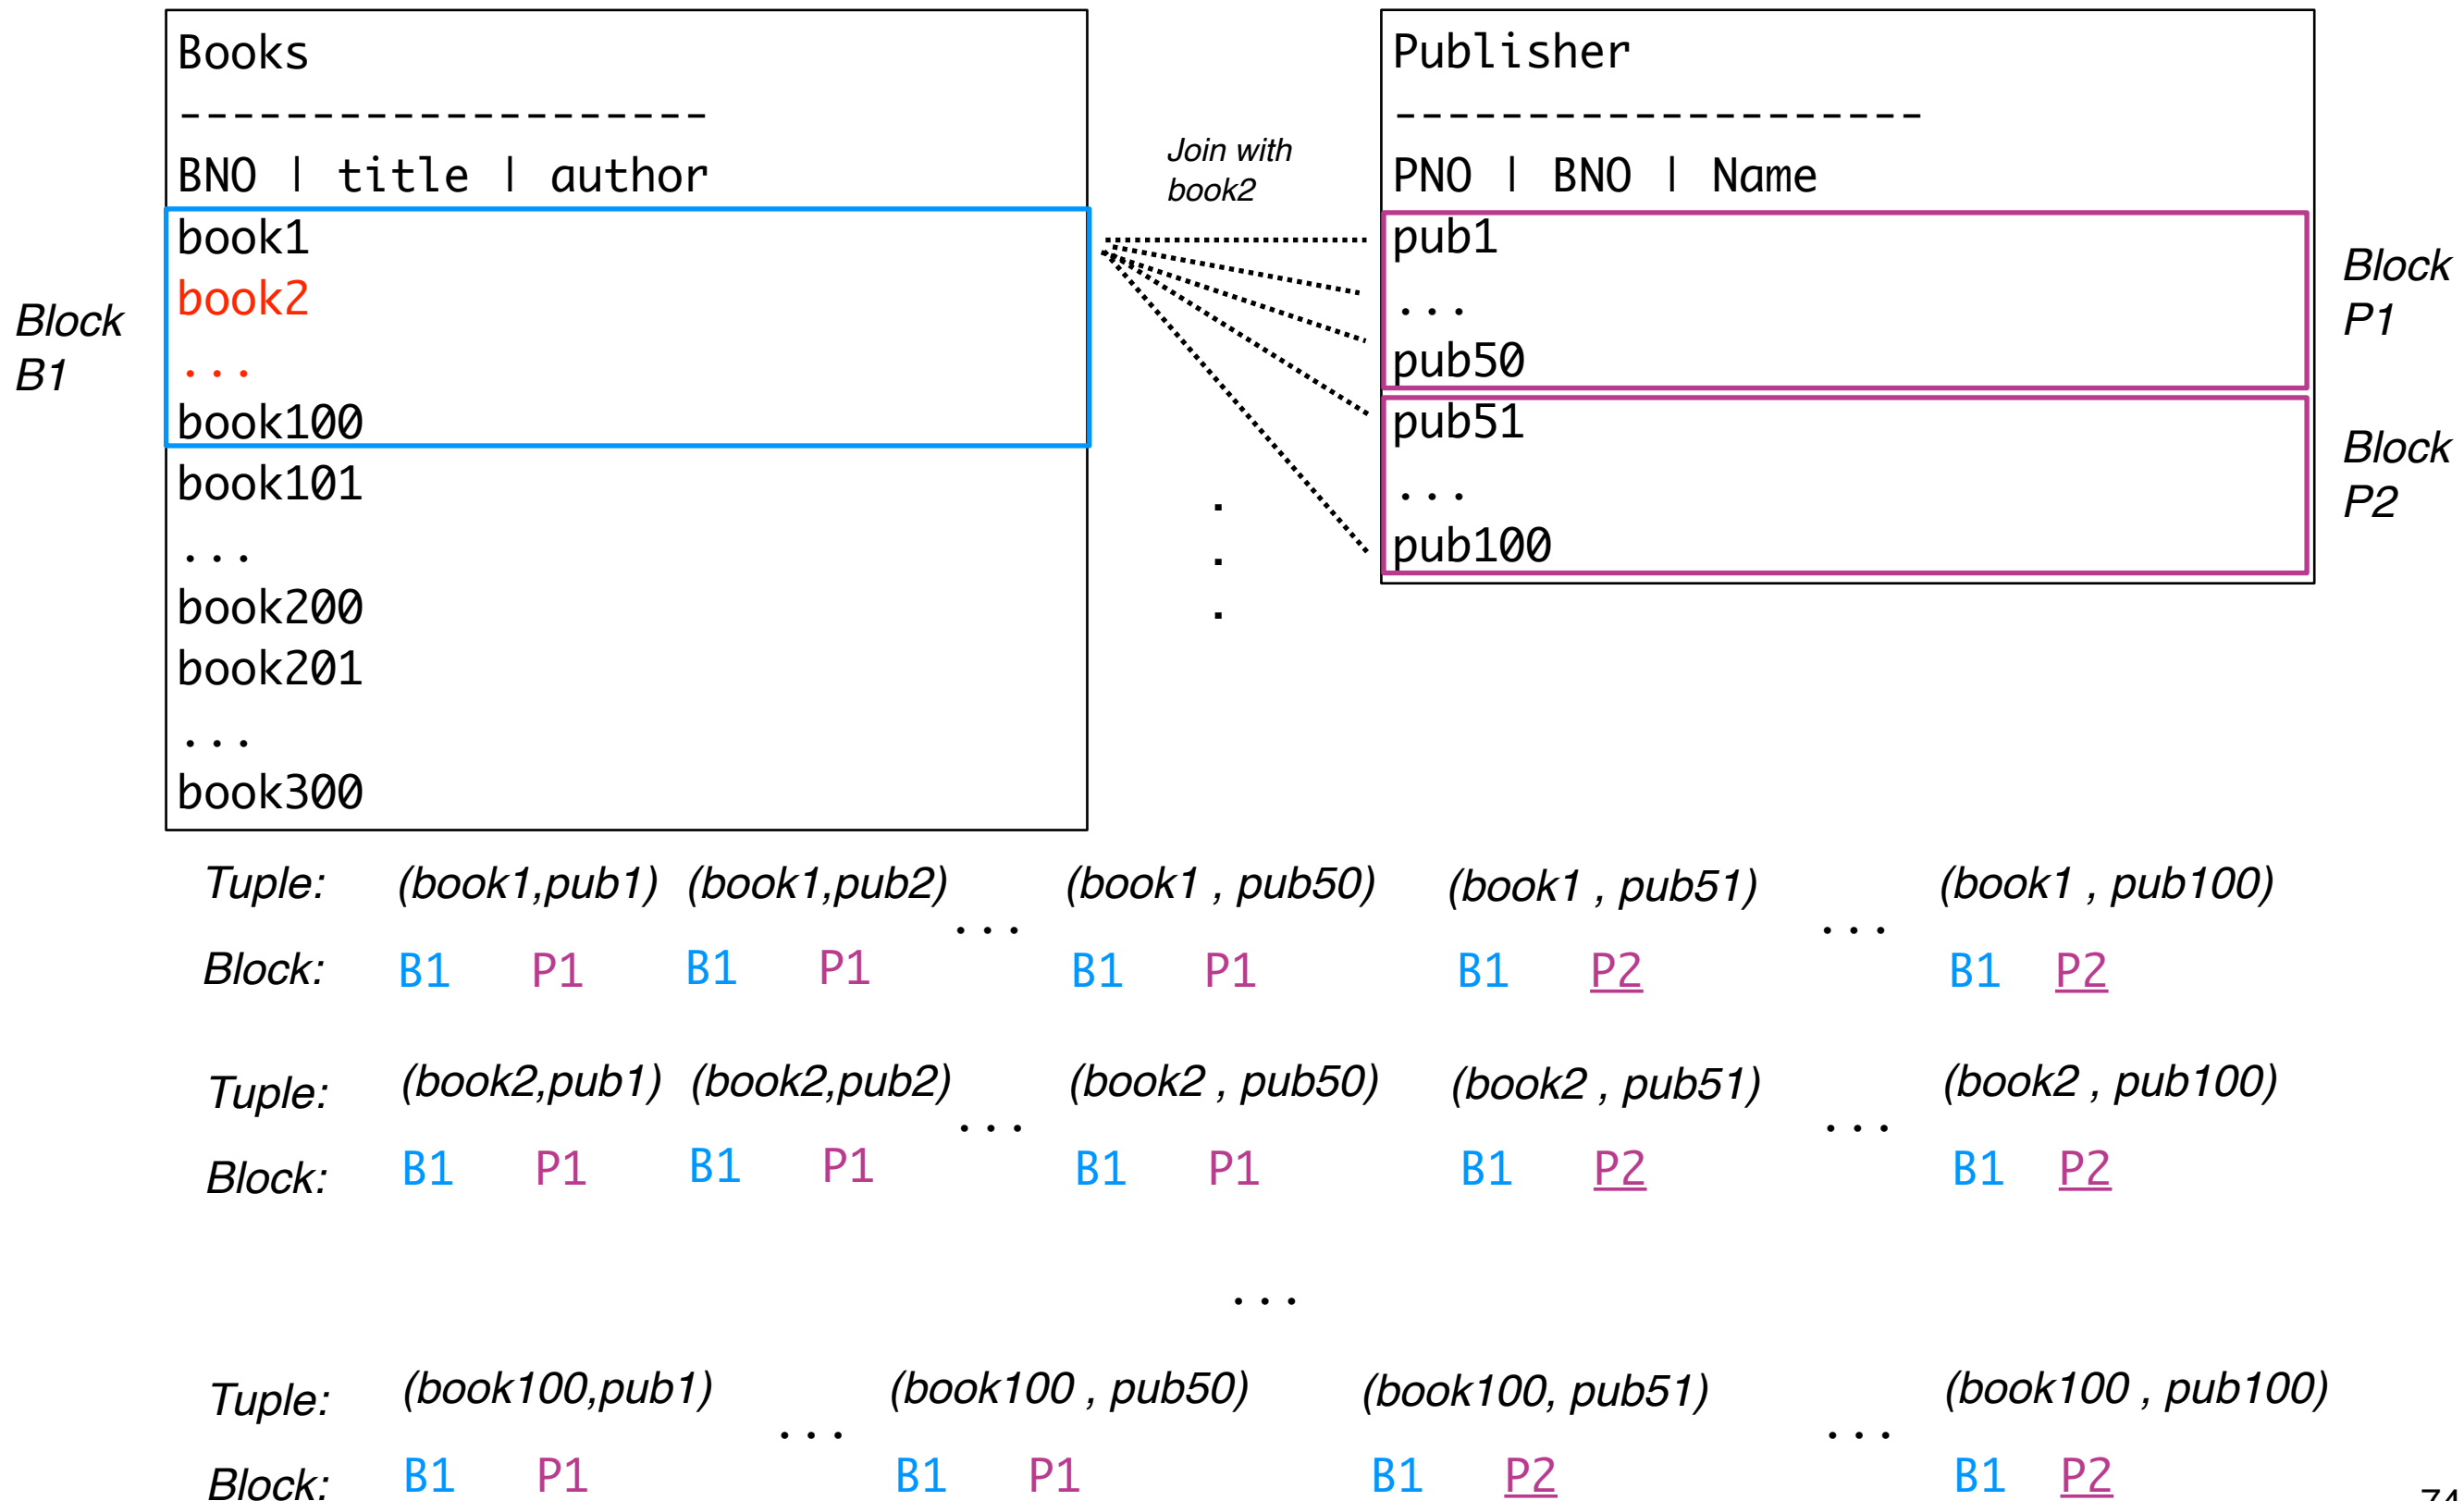


Number of misses = 3 (same as FIFO for this query)

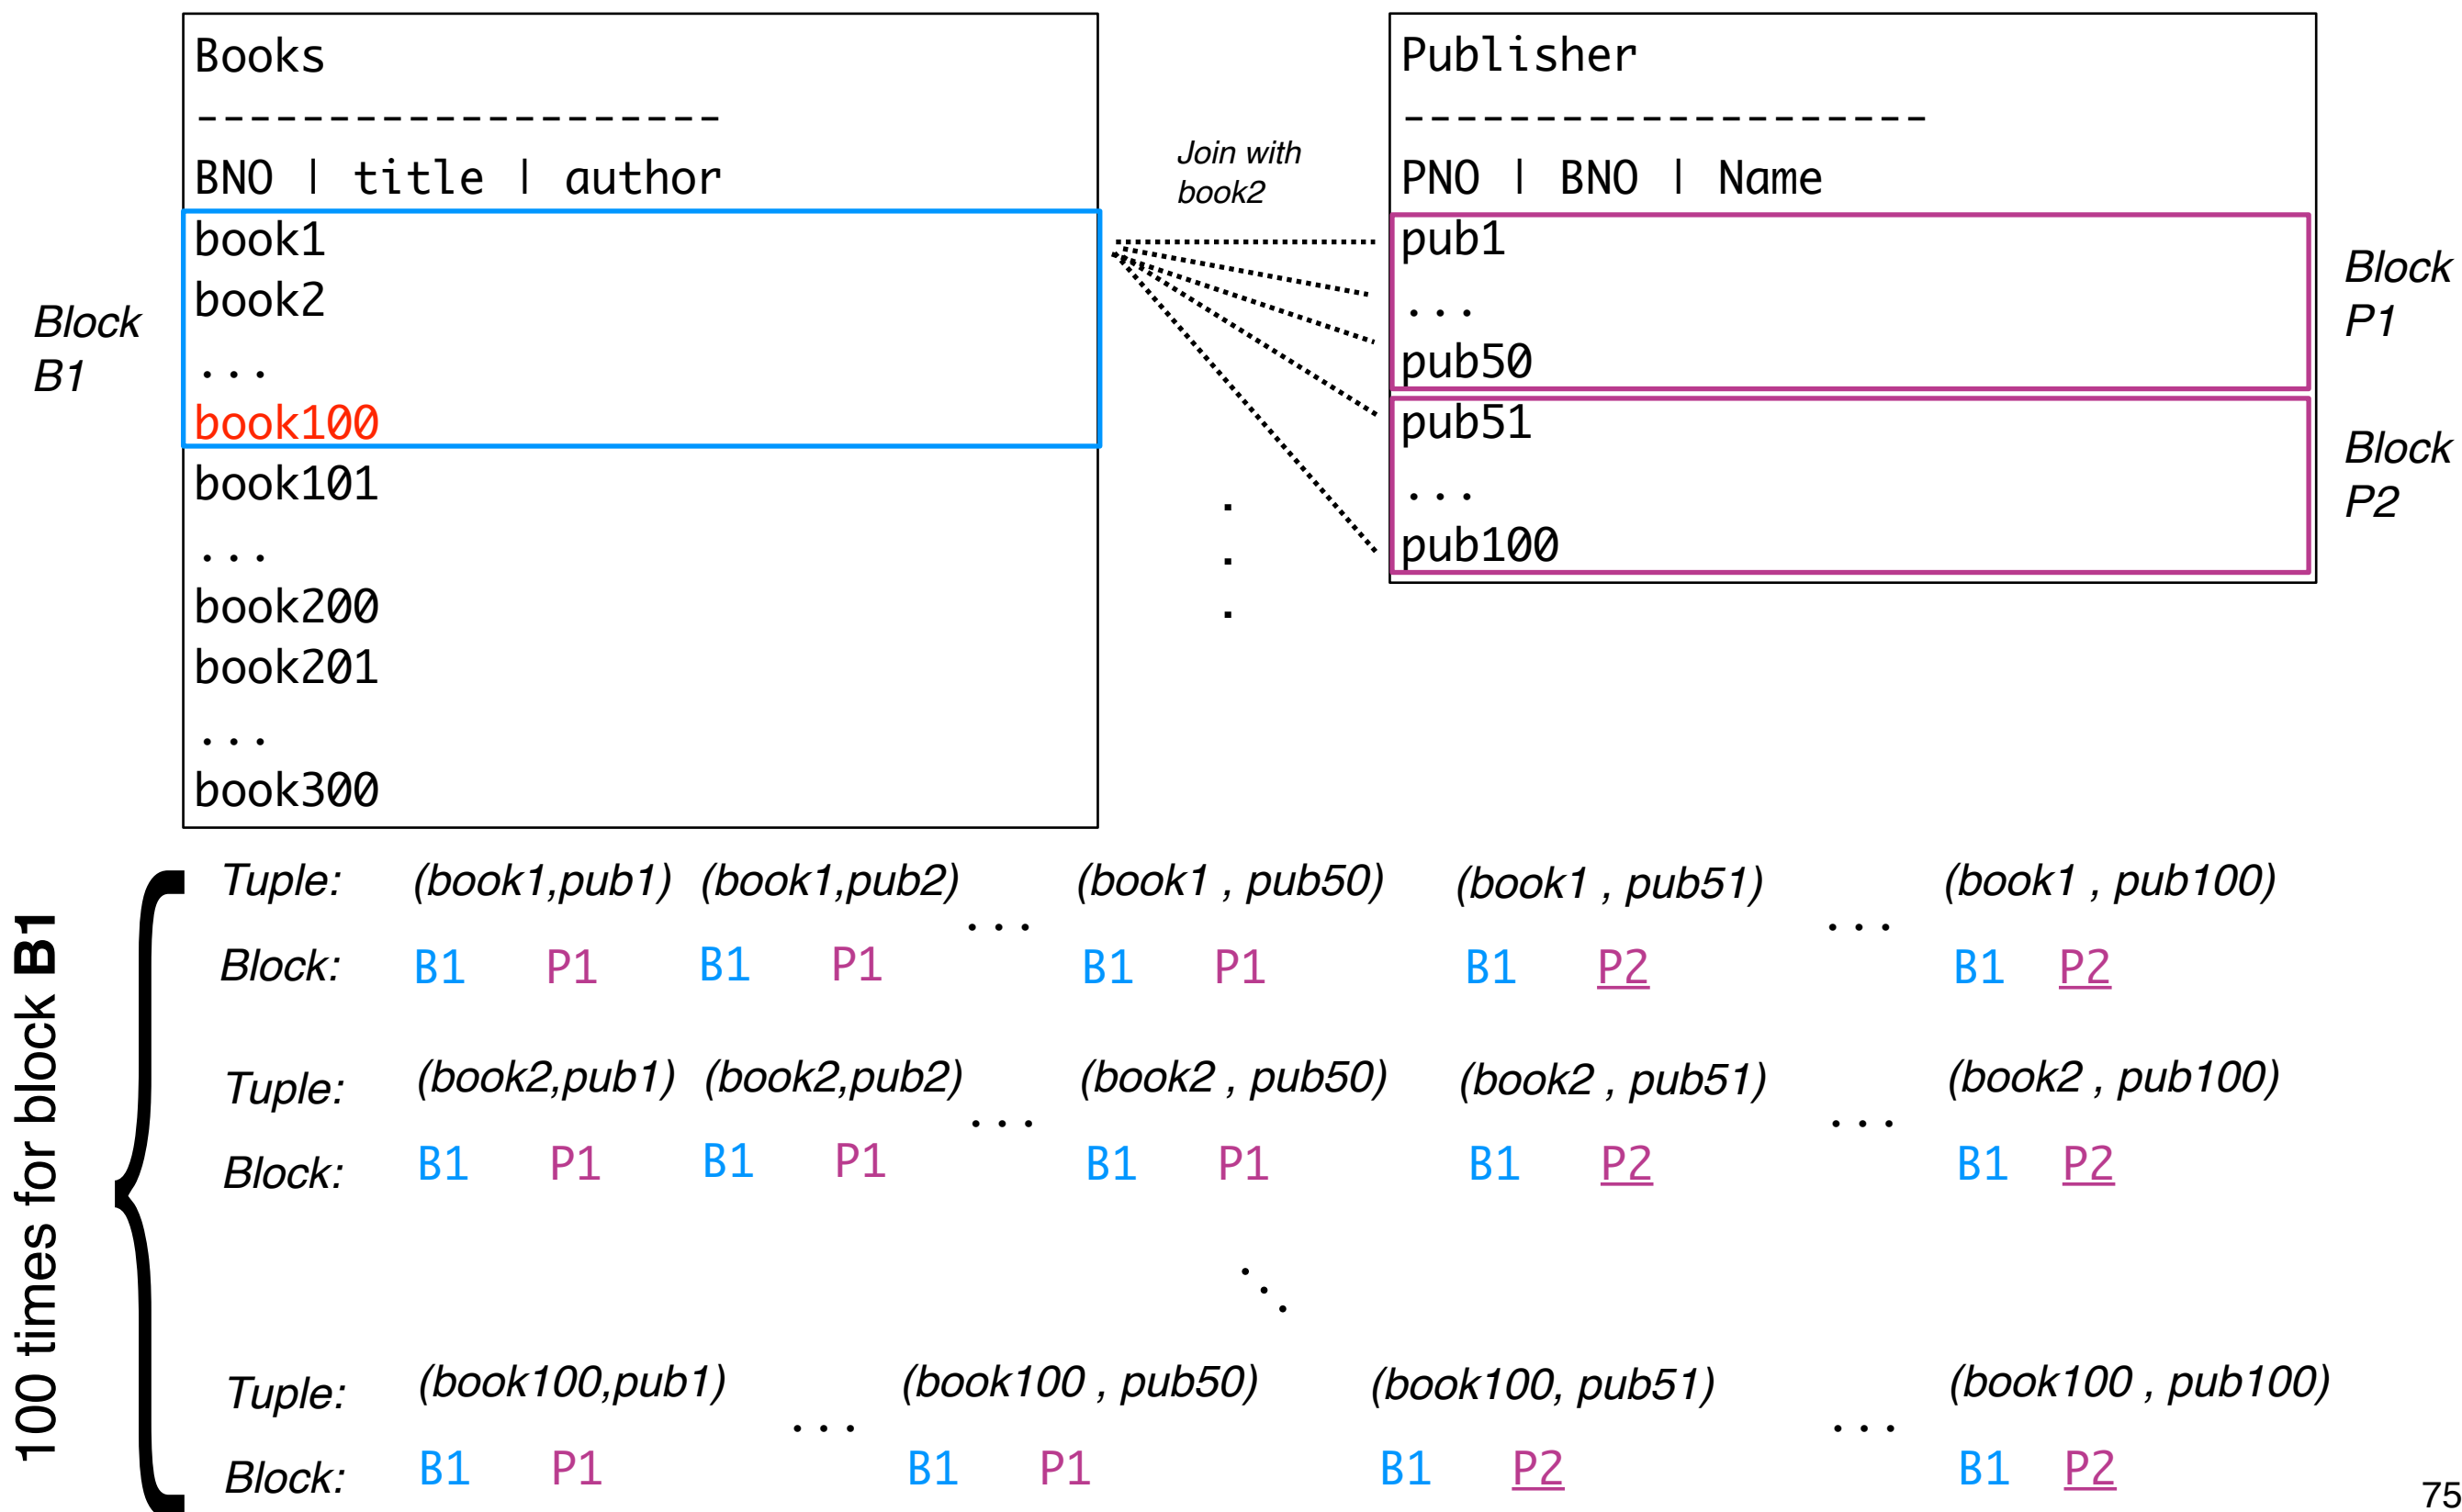
Access Pattern for a Join



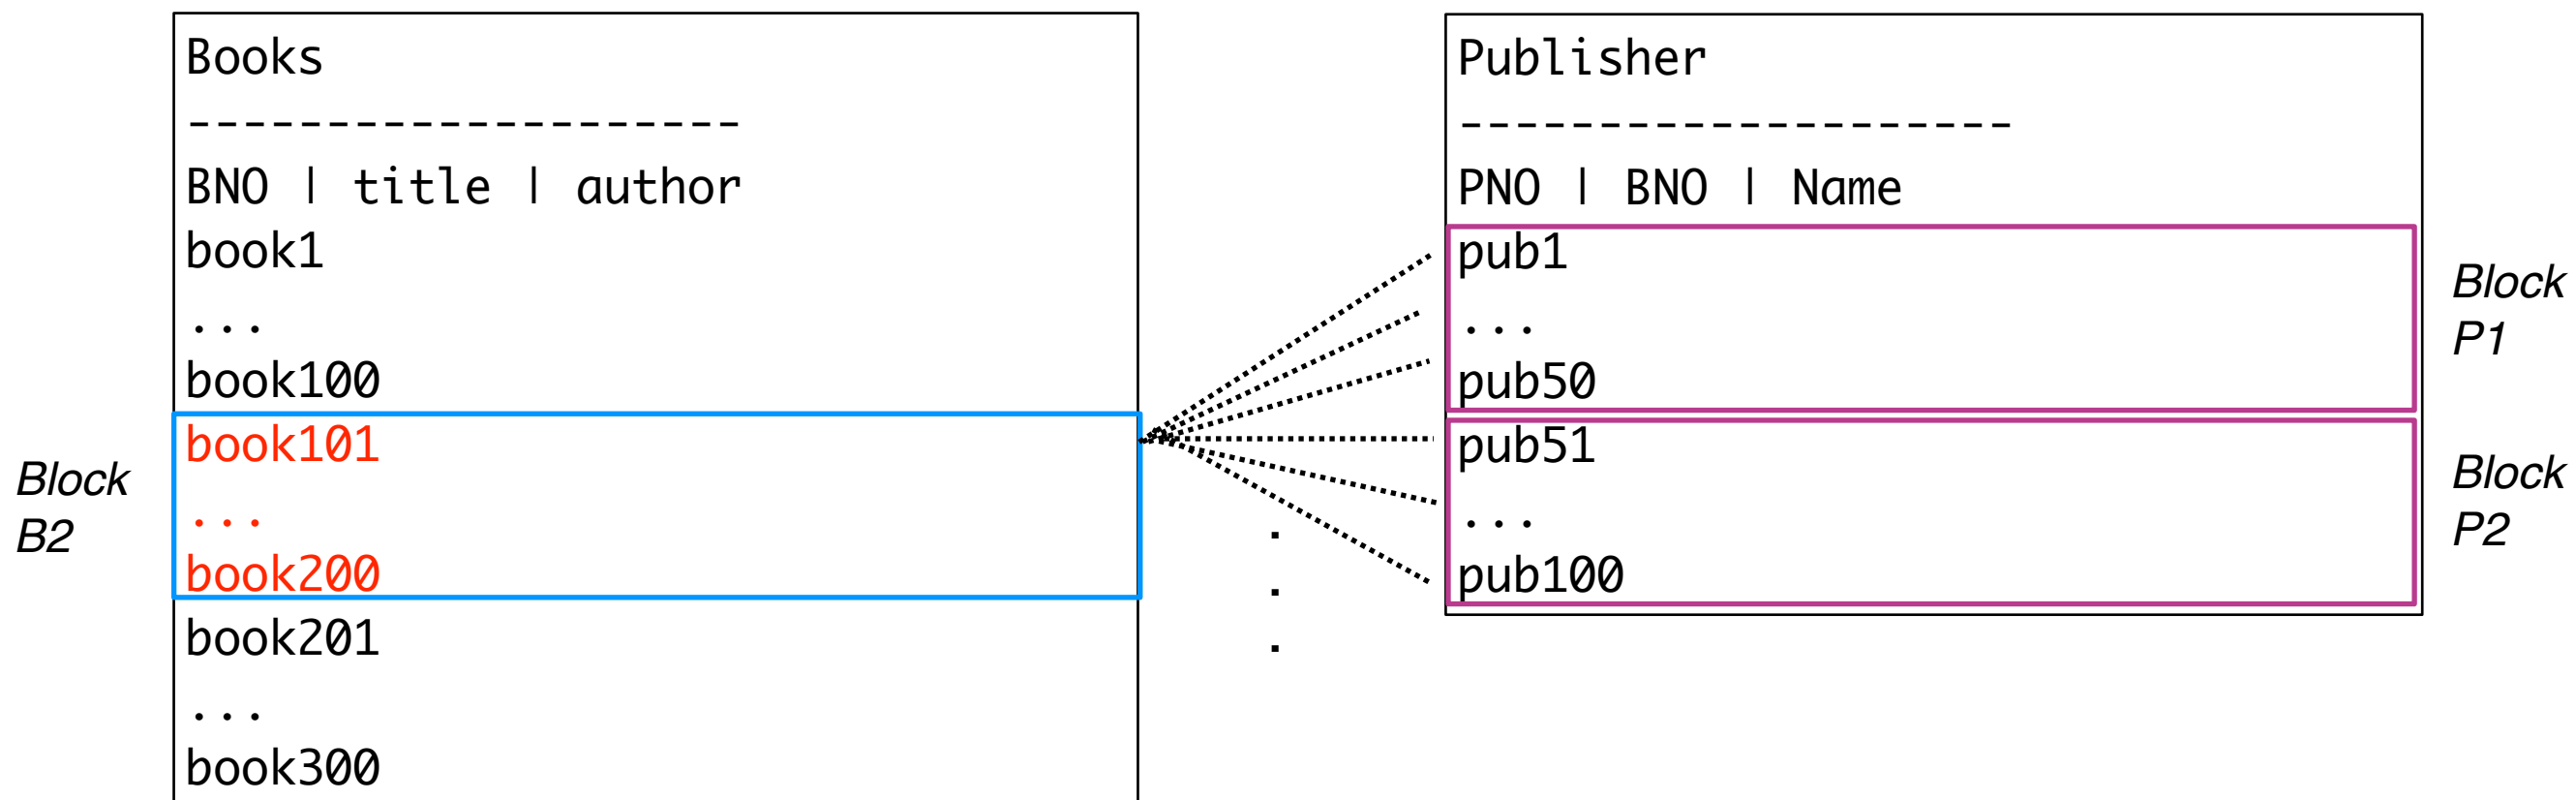
Access Pattern for a Join



Access Pattern for a Join



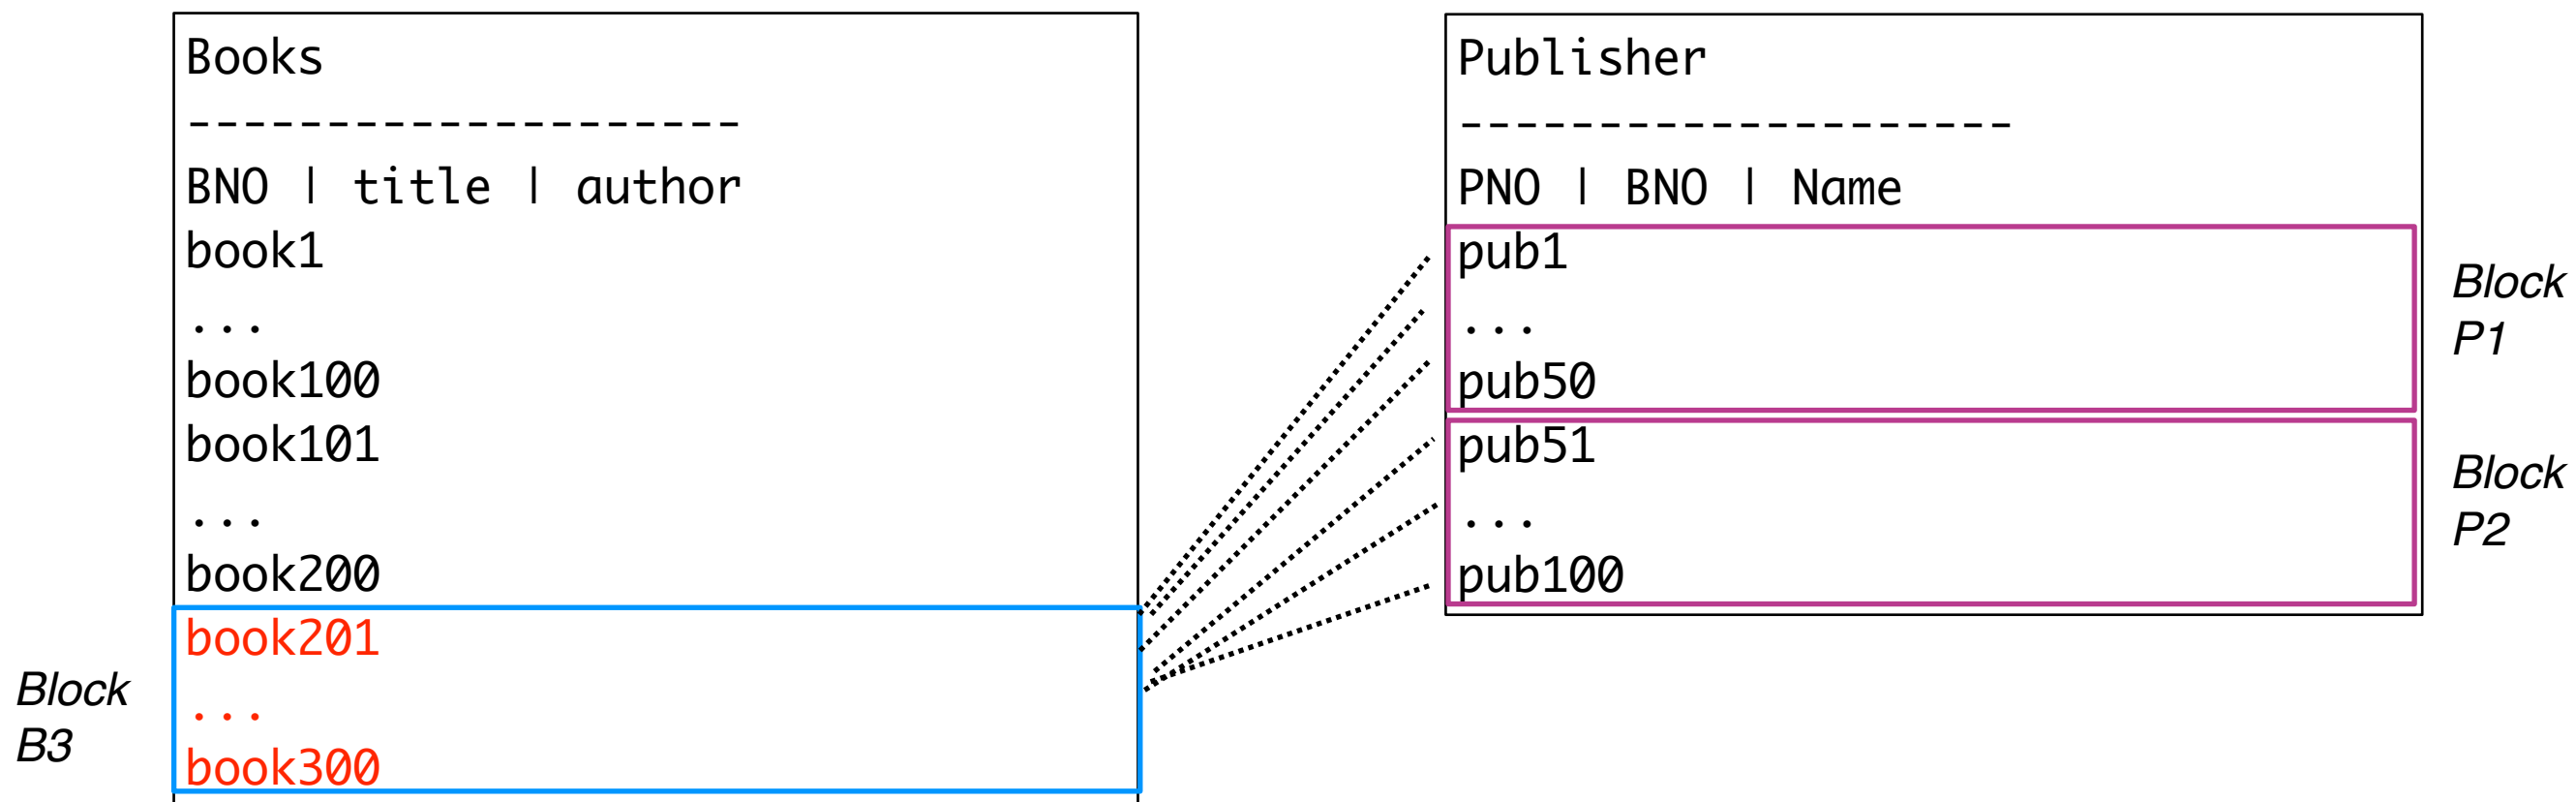
Access Pattern for a Join (Cont.)



100 times for block **B2**

<i>Tuple:</i>	$(b101, p1)$	$(b101, p2)$...	$(b101, p50)$	$(b101, p51)$	$(b101, p52)$...	$(b101, p100)$
<i>Block:</i>	B2	P1	B2	P1	B2	P2	B2	P2
<i>Tuple:</i>	$(b102, p1)$	$(b102, p2)$...	$(b102, p50)$	$(b102, p51)$	$(b102, p52)$...	$(b102, p100)$
<i>Block:</i>	B2	P1	B2	P1	B2	P2	B2	P2
...								
<i>Tuple:</i>	$(b200, p1)$	$(b200, p2)$...	$(b200, p50)$	$(b200, p51)$	$(b200, p52)$...	$(b200, p100)$
<i>Block:</i>	B2	P1	B2	P1	B2	P2	B2	P2

Access Pattern for a Join (Cont.)




100 times for block B3

Tuple:	(b201 , p1)	(b201 , p2)	...	(b201 , p50)	(b201 , p51)	(b201 , p52)	...	(b201 , p100)
Block:	B3 P1	B3 P1	...	B3 P1	B3 P2	B3 P2	...	B3 P2
Tuple:	(b202 , p1)	(b202 , p2)	...	(b202 , p50)	(b202 , p51)	(b202 , p52)	...	(b202 , p100)
Block:	B3 P1	B3 P1	...	B3 P1	B3 P2	B3 P2	...	B3 P2
...								
Tuple:	(b300 , p1)	(b300 , p2)	...	(b300 , p50)	(b300 , p51)	(b300 , p52)	...	(b300 , p100)
Block:	B3 P1	B3 P1	...	B3 P1	B3 P2	B3 P2	...	B3 P2

FIFO Replacement Policy

- Database buffer allocates **two** blocks in RAM
 - Access Pattern of **natural join**(Book, Publisher)



2 misses in first 50 joins:
Must fetch **B1**
and **P2** from
disk initially.


2 misses
This completes the
join of tuple
book1
with all tuples in
Publishers

1 miss
Next 50 joins
for **book2**

2 misses
Remaining 50 joins
for **book2**

LRU Replacement Policy

- Database buffer allocates **two** blocks in RAM
 - Access Pattern of **natural join**(Book, Publisher)



2 misses in first 50 joins:
Must fetch **B1**
and **P2** from
disk initially.

1 miss
This completes the
join of tuple
book1
with all tuples in
Publishers

1 miss
Next 50 joins
for **book2**

1 miss
Remaining 50 joins
for **book2**

Evaluation of FIFO vs. LRU for a Join

► Full Access pattern

- (DAT = disk access time per block; MAT = RAM access time per block)

Full Access Pattern for NLJ	No-Buffering Misses	FIFO Misses	LRU Misses
B1, P1 (repeated 50 times)	50 +	2 +	2 +
B1, P2 (repeated 50 times)	50 +	3 +	2 +
...	50 +	3 +	2 +
B1, P1 (repeated 50 times)
B1, P2 (repeated 50 times)	50 = 5000 misses	3 = 150 misses	2 = 100 misses
B2, P1 (repeated 50 times)	50 +	2 +	2 +
B2, P2 (repeated 50 times)	50 +	3 +	2 +
...	50 +	3 +	2 +
B2, P1 (repeated 50 times)
B2, P2 (repeated 50 times)	50 = 5000 misses	3 = 150 misses	2 = 100 misses
B3, P1 (repeated 50 times)	50 +	2 +	2 +
B3, P2 (repeated 50 times)	50 +	3 +	2 +
...	50 +	3 +	2 +
B3, P1 (repeated 50 times)
B3, P2 (repeated 50 times)	50 = 5000 misses	3 = 150 misses	2 = 100 misses
Total Misses for a Join	15000 misses	450 misses	300 misses
Total Query Time (Assume DAT = 10ms MAT = 0 ms)	150000 ms (= 2.5 min)	4500 ms (= 4.5 sec)	3000 ms (= 3 sec)

Topics

- ▶ Disks
 - Performance Metrics
- ▶ Database File Structure
- ▶ Tuple Organization in Files
- ▶ Buffer Manager
- ▶ Conclusion



Administrivia 10/29

► Happy Halloween!

► Some **spooky** reminders:

- Project proposal due Monday 11/1
- Homework 5 due Friday 11/5

► Last time...

• **Files => Blocks => Disk Sectors**

- Importance of blocks arranged contiguously on disk
- Importance of file access pattern (sequential vs. irregular)

► Today

- Modeling data access time
- Disk scheduling



Administrivia 11/1

- ▶ Reminders:
 - Project proposals due tonight!
 - Hwk 5 due Friday
- ▶ Last time...
 - Modeling disk access time
- ▶ Today
 - Disk scheduling policies
 - Buffer/cache management